

# CSCI 8530

# Advanced Operating Systems

Part 2

Organization of an Operating System

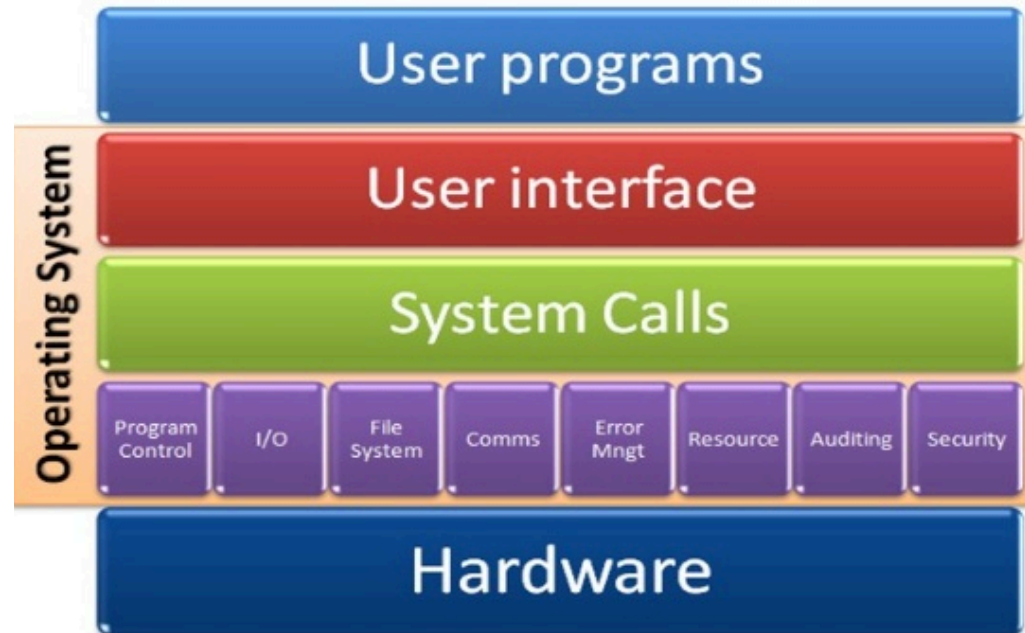
# Review: What Is An Operating System?

- Provides abstract computing environment
- Supplies computational services
- Manages resources
- Hides low-level hardware details
- Note: operating system software is among the most complex ever devised



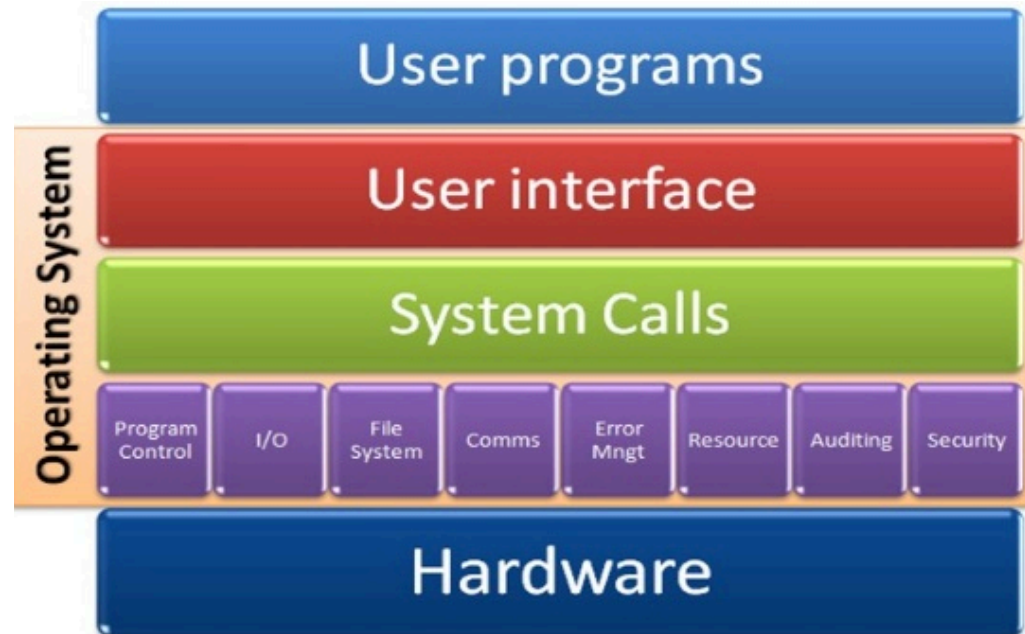
# Review: Example Services an OS Supplies

- Support for concurrent execution
- Process synchronization
- Inter-process communication mechanisms
- Message passing and asynchronous events
- Management of address spaces and virtual memory
- Protection among users and running applications
- High-level interface for I/O devices
- A file system and file access facilities
- Intermachine communication



# Review: What an Operating System is NOT

- Hardware
- Language
- Compiler
- Windowing system or browser
- Command interpreter
- Library of utility functions
- Graphical desktop

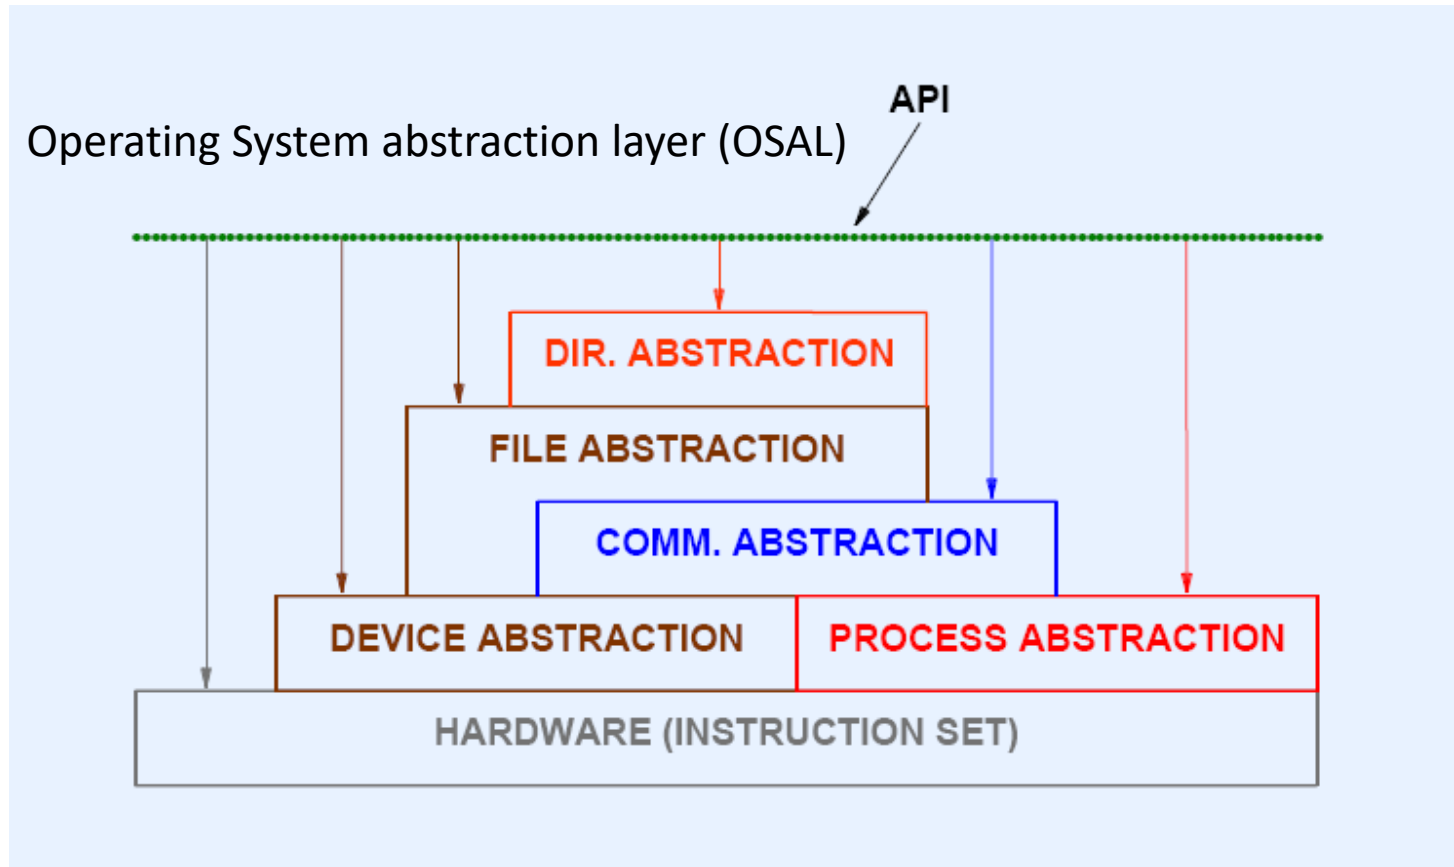


# **AN OPERATING SYSTEM FROM THE OUTSIDE**

# The System Interface

- Single copy of OS per computer
  - Hidden from users
  - Accessible only to application programs
- *Application Program Interface (API)*
  - Defines services OS makes available
  - Defines parameters for the services
  - Provides access to all abstractions
  - Hides hardware details

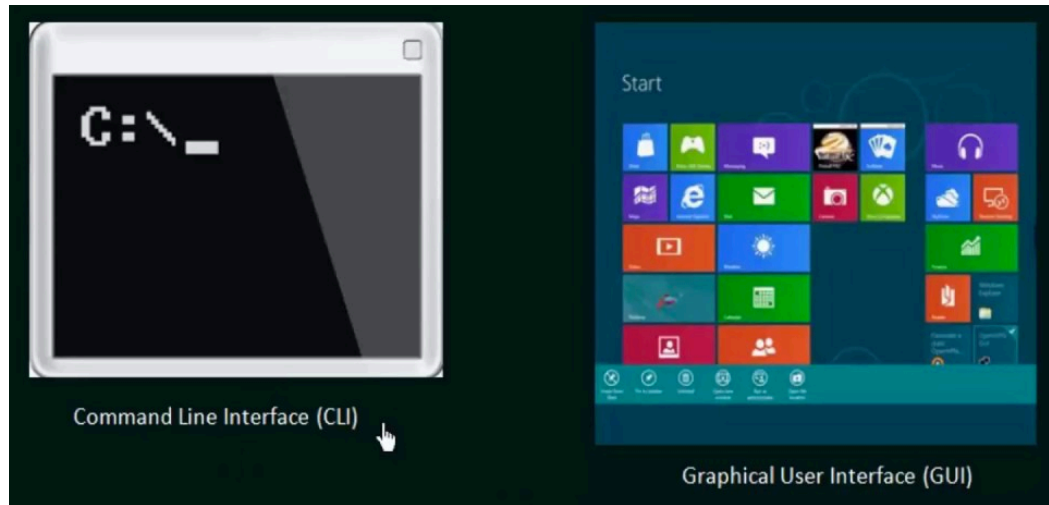
# OS Abstractions and Application Interface



- Modules in the OS offer services
- Some services build on others

# Interface to System Services

- Appears to operate like a function call mechanism
  - OS makes set of “functions” available to applications
  - Application supplies arguments using standard mechanism
  - Application “calls” one of the OS functions
- Control transfers to OS code that implements the function
- Control returns to caller when function completes



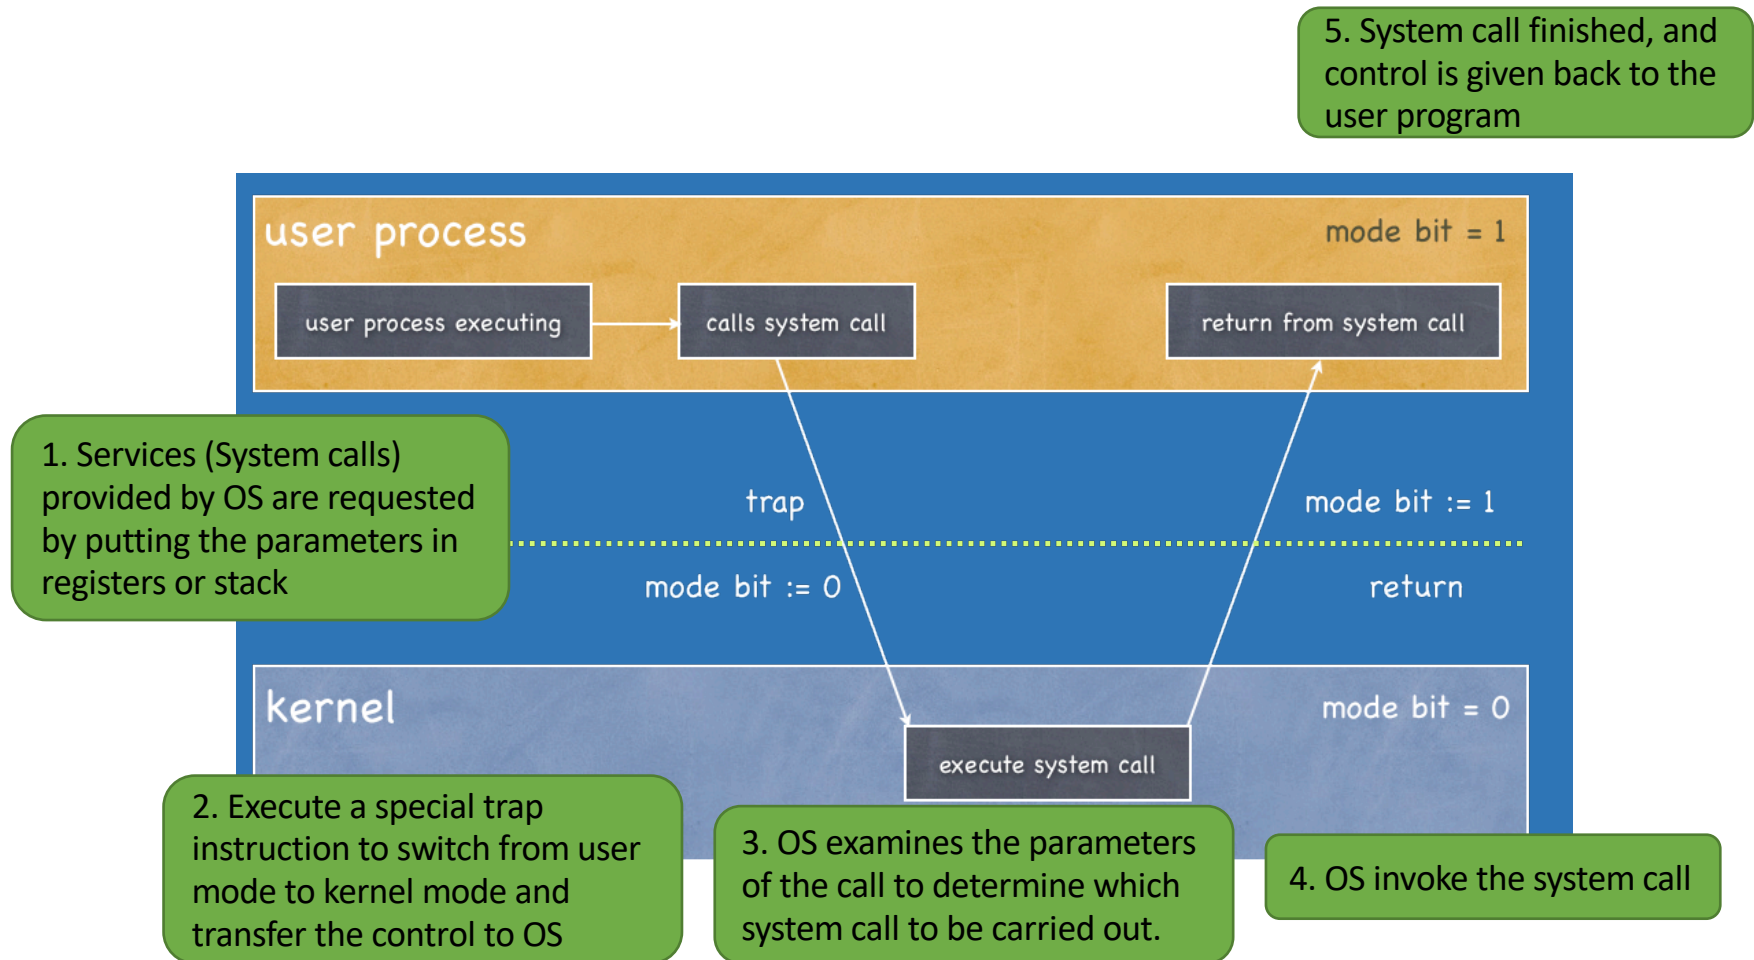


# Interface to System Services

## (continued)

- Requires special instruction to invoke OS function
  - Moves from application *address space* to OS
  - Changes from application *mode* or *privilege level* to OS
- Terminology used by various vendors
  - *System call*
  - *Trap/exception*
  - *Supervisor call*
- We will use the generic term *system call*
  - *A system call*
    - The programmatic way in which a computer program requests a service from the kernel of OS it is executed on
    - This way for programs to interact with OS

# System Calls: A Closer Look



# Example System Call in Xinu: Write a Character on the Console

```
/* ex1.c - main */
```

```
#include <xinu.h>
```

```
/*-----  
 * main - write "hi" on the console  
 *-----  
*/
```

```
void main(void)
```

```
{
```

```
    putc(CONSOLE, 'h');
```

```
    putc(CONSOLE, 'i');
```

```
    putc(CONSOLE, '\n');
```

```
}
```

# OS Services and System Calls

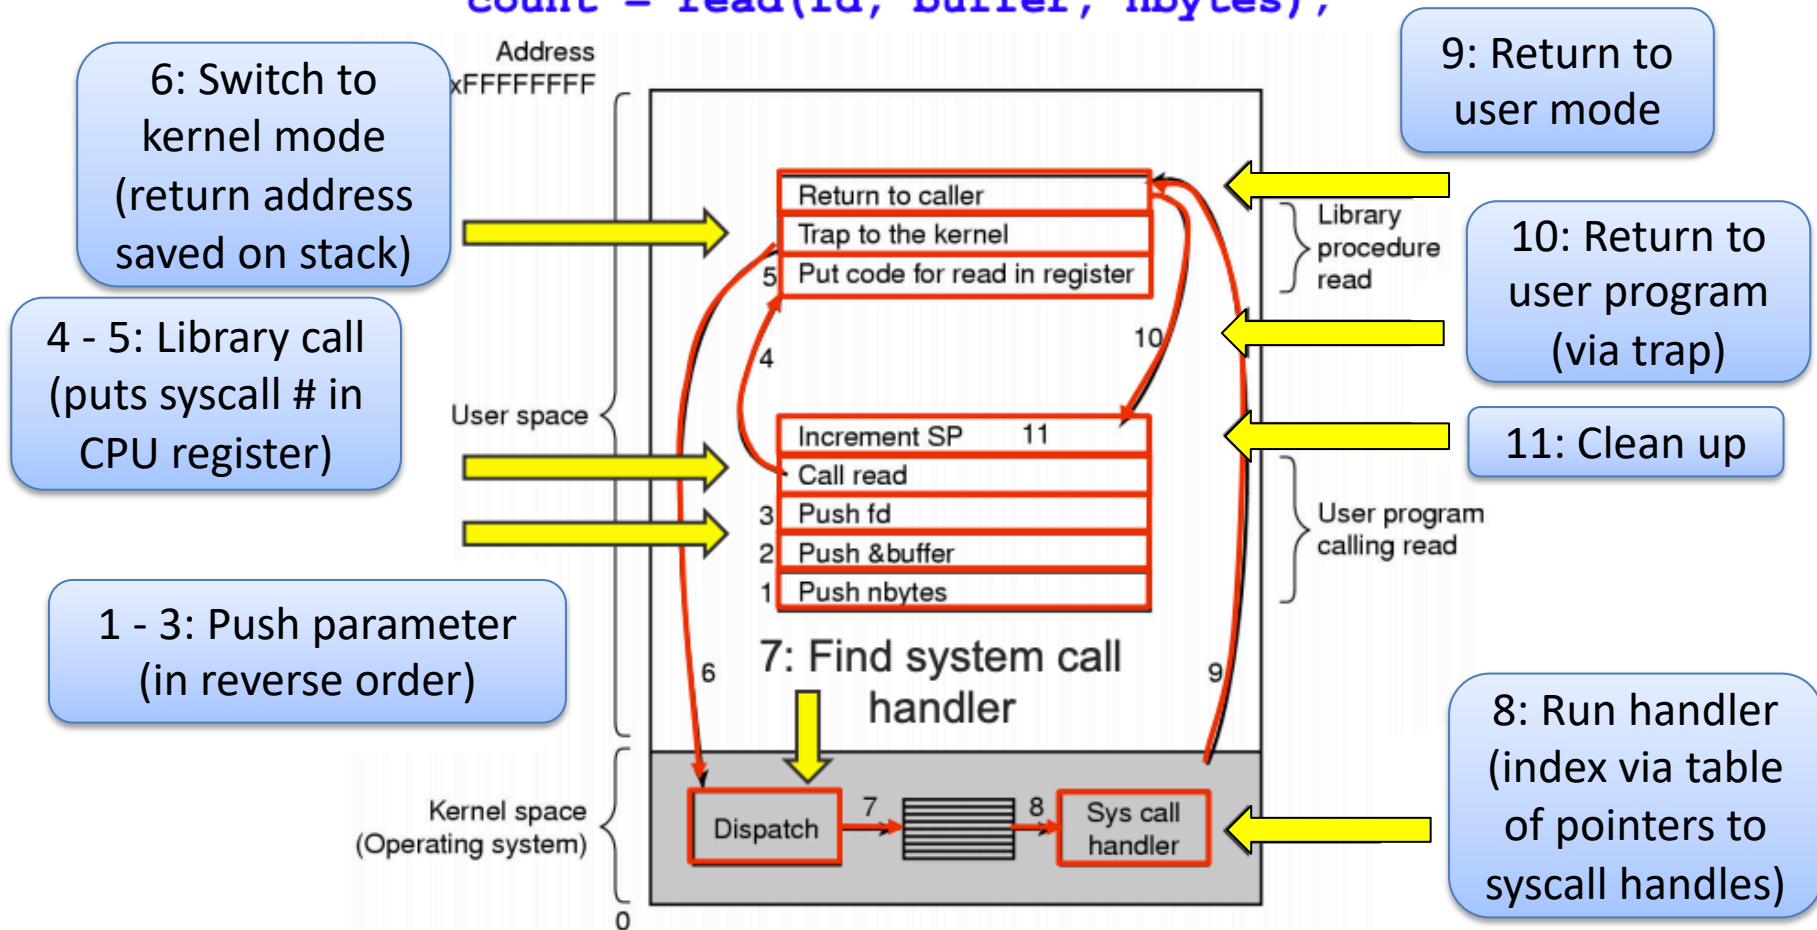
- OS provides services accessed through system call interface
- Most services employ a set of several system calls
- Examples
  - Process management service includes functions: *suspend* and then *resume* a process
  - Socket API used for Internet communication includes many functions: establish a socket on the *client* side
    - Create a socket with the *socket()* system call
    - Connect the socket to the address of the server using the *connect()* system call
    - Send and receive data (used with I/O)

# System Calls Used With I/O

- An I/O subsystem comprises of I/O devices and their corresponding driver software
  - Drivers hide the peculiarities of specific hardware devices from the users.
- Open-close-read-write paradigm
  - Application
    - Uses *open* to connect to a file or device
    - Calls functions to *write* data or *read* data
    - Calls *close* to terminate use
  - Internally, the set of I /O functions coordinate
    - *Open* returns a descriptor
    - *Read* and *write* operate on descriptor

# Steps for Making a System Call (Example: read call)

```
count = read(fd, buffer, nbytes);
```



# Implementing a System Call

- System calls are often implemented using traps
  - OS gains control through trap
  - Switches to supervisor mode
  - Performs the service
  - Switches back to user mode
  - Gives control back to user

```
movl $1, %eax  
int $0x80
```

## Which call?

1: exit  
2: fork  
3: read  
4: write  
5: open  
6: close  
...

Trap to the OS

System-call specific arguments are put in registers

# Major System Calls (1/2)

Process Management	
<code>pid = fork( )</code>	Create a child process identical to the parent
<code>pid = waitpid(pid, &amp;statloc, options)</code>	Wait for a child to terminate
<code>s = execve(name, argv, environp)</code>	Replace a process' core image
<code>exit(status)</code>	Terminate process execution and return status

File Management	
<code>fd = open(file, how, ...)</code>	Open a file for reading, writing or both
<code>s = close(fd)</code>	Close an open file
<code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
<code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
<code>position = lseek(fd, offset, whence)</code>	Move the file pointer
<code>s = stat(name, &amp;buf)</code>	Get a file's status information



# Major System Calls (2/2)

Directory and File System Management	
<code>s = mkdir(name, mode)</code>	Create a new directory
<code>s = rmdir(name)</code>	Remove an empty directory
<code>s = link(name, name)</code>	Create a new entry, name, pointing to name
<code>s = unlink(name)</code>	Remove a directory entry
<code>s = mount(special, name, flag)</code>	Mount a file system
<code>s = umount(special)</code>	Unmount a file system

Miscellaneous	
<code>s = chdir(dirname)</code>	Change the working directory
<code>s = chmod(name, mode)</code>	Change a file's protection bits
<code>s = kill(pid, signal)</code>	Send a signal to a process
<code>seconds = time(&amp;seconds)</code>	Get the elapsed time since January 1, 1970

# Question

- What kinds of system calls does an OS need?
  - Process creation and management
  - Main memory management
  - File Access, Directory and File system management
  - Device handling(I/O)
  - Protection
  - Networking, etc.



# Question

Why does the OS control I/O?


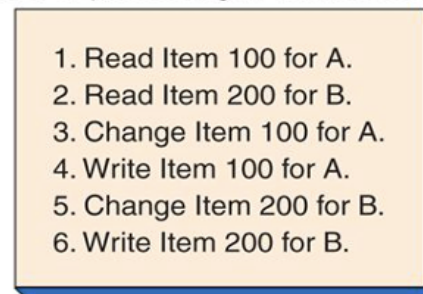
- Safety
  - The computer must ensure that if a program has a bug in it, then it doesn't crash or mess up
    - The system
    - Other programs that may be running at the same time or later
- Fairness
  - Make sure other programs have a fair use of device

# Concurrent Processing

- Fundamental concept dominates OS design
- *Real concurrency* achieved by hardware
  - I/O devices operate at same time as processor
  - Multiple processors/cores each operate at the same time
- *Apparent concurrency* achieved with multitasking (multiprogramming)
  - Multiple programs appear to operate simultaneously
  - OS provides the illusion
  - Example: User A and User B

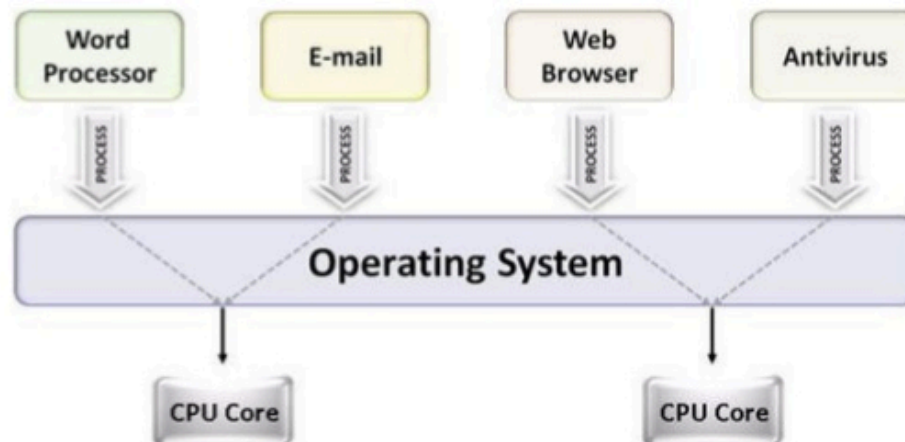


Order of processing at database server



# Multitasking

- Powerful abstraction
- Allows user(s) to run multiple computations
- OS switches processor(s) among available computations quickly
- All computations appear to proceed in parallel
- Scheduler

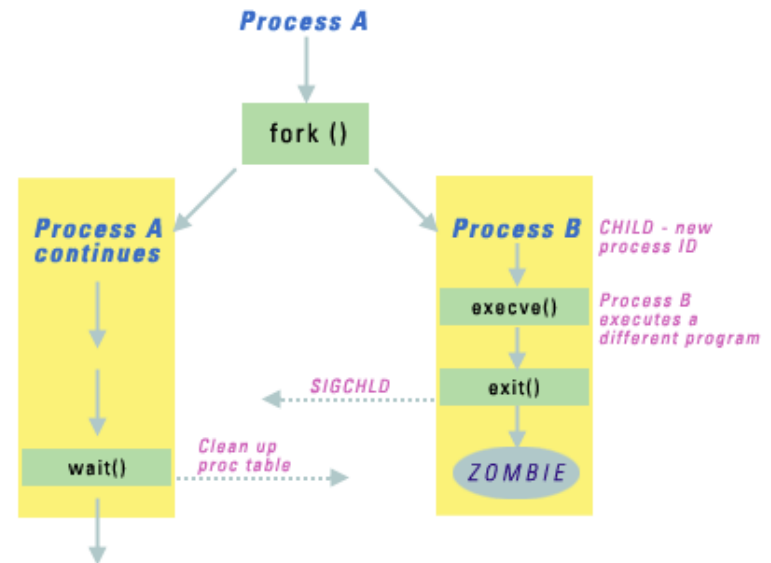


# Terminology Used with Multitasking

- *Program* consists of static code and data
- *Function* is a unit of application program code
- *Process* (also called a *thread of execution*) is an active computation (i.e., the execution or “running” of a program)

# Process

- A process is an instance of a running program
- Managed entirely by OS; unknown to hardware
- Process provides each program with two key OS abstractions
  - **Logical control flow**
    - Each program seems to have exclusive use of the CPU
  - **Private virtual address space**
    - Each program seems to have exclusive use of main memory
- How are these illusions maintained?
  - Process executions interleaved (multitasking) or run on separate cores
  - Address spaces managed by virtual memory system
- Operates *concurrently* with other processes



# Example of Process Creation in Xinu

```
/* ex2.c - main, sndA, sndB */
```

```
#include <xinu.h>
```

```
void sndA(void), sndB(void);
```

```
/*-----  
 * main - Example of creating processes in Xinu  
 *-----  
 */
```

```
void main(void)
```

```
{  
    resume( create(sndA, 1024, 20, "process 1", 0) );  
    resume( create(sndB, 1024, 20, "process 2", 0) );  
}
```

```
/*-----  
 * sndA - Repeatedly emit 'A' on the console without terminating  
 *-----  
 */
```

```
void sndA(void)
```

```
{  
    while( 1 )  
        putc(CONSOLE, 'A');  
}
```




# Example of Process Creation in Xinu

```
/*-----  
 * sndB - Repeatedly emit 'B' on the console without terminating  
 *-----  
 */  
void sndB(void)  
{  
    while( 1 )  
        putc(CONSOLE, 'B');  
}
```

# Question


- System Calls V.S. Function Calls?

System Calls



- OS is trusted; user is not.
- OS has super-privileges; user does not
- Must take measures to prevent abuse

Function Calls



Caller and callee are in the same Process

- Same user
- Same "domain of trust"

# Xinu Difference Between Function Call and Process Creation

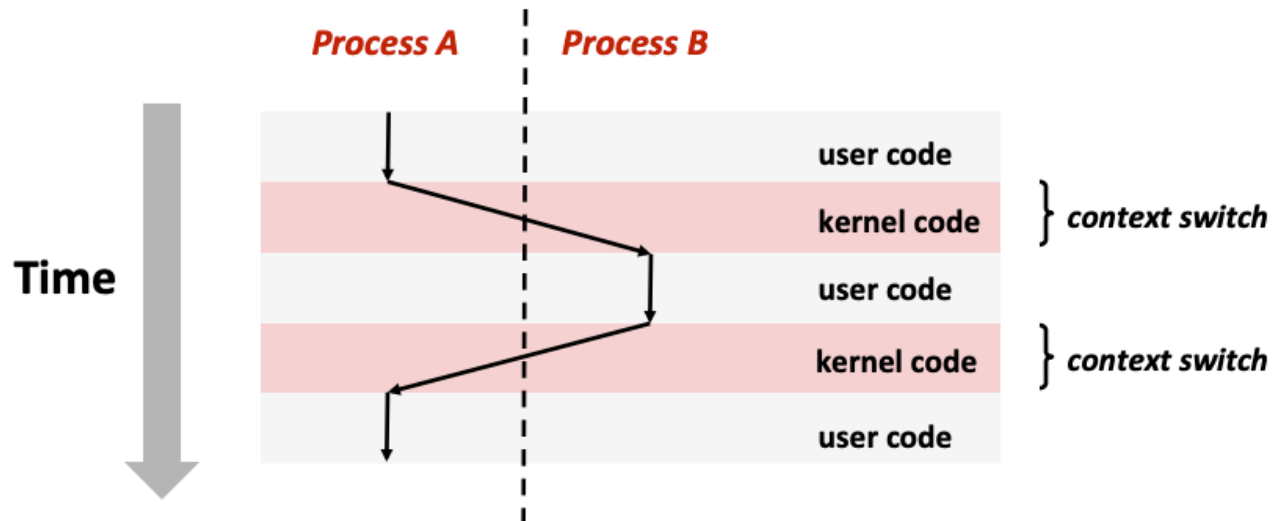
- Normal function call
  - Synchronous execution
  - Single computation
- The “*create*” system call that starts a new process
  - Asynchronous execution
  - Two processes proceed after the call

# Distinction Between a Program and a Process

- Sequential program
  - Set of functions executed by a single thread of control
- Process
  - Computational abstraction not usually part of the programming language
  - Created independent of code that is executed
  - Key idea: multiple processes can execute the same code concurrently
- In the following example, two processes execute function *sndch* concurrently

# Context Switching

- Processes are managed by a shared chunk of OS code called the *kernel*
  - Important: the kernel is not a separate process, but rather runs as part of some user process
- Control flow passes from one process to another via a context switch



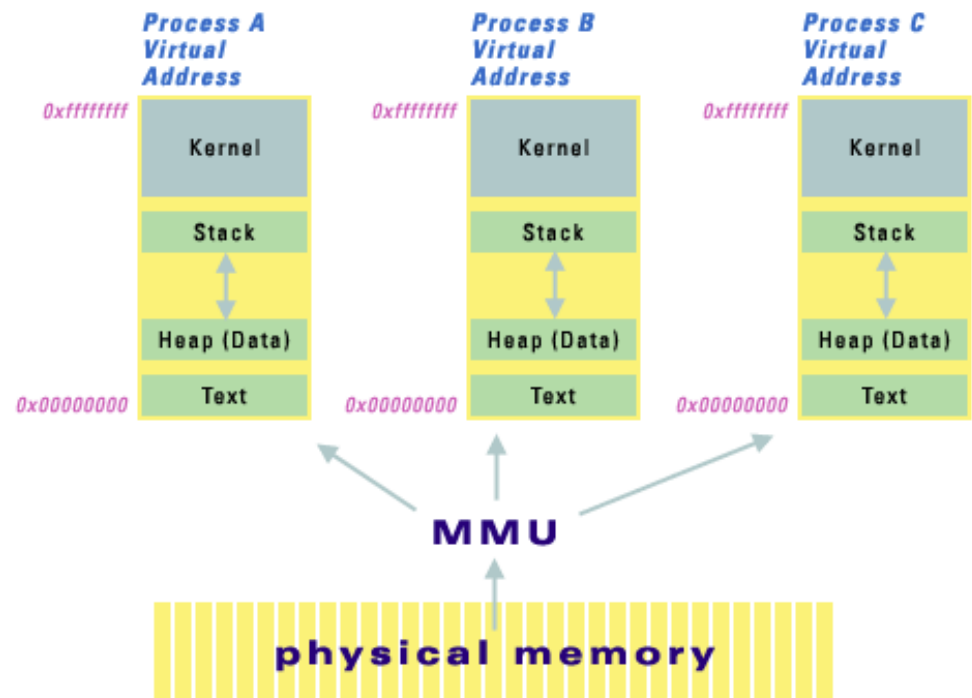
# Example of Two Processes Sharing Code

```
/* ex3.c - main, sndch */
#include <xinu.h>
void sndch(char);
/*-----
 * main - Example of 2 processes executing the same code concurrently
 *-----
 */
void main(void)
{
    resume( create(sndch, 1024, 20, "send A", 1, 'A') );
    resume( create(sndch, 1024, 20, "send B", 1, 'B') );
}

/*-----
 * sndch - Output a character on a serial device indefinitely
 *-----
 */
void sndch(
    char ch          /* The character to emit continuously */
)
{
    while ( 1 )
        putc(CONSOLE, ch);
}
```

# Storage Allocation When Multiple Processes Execute

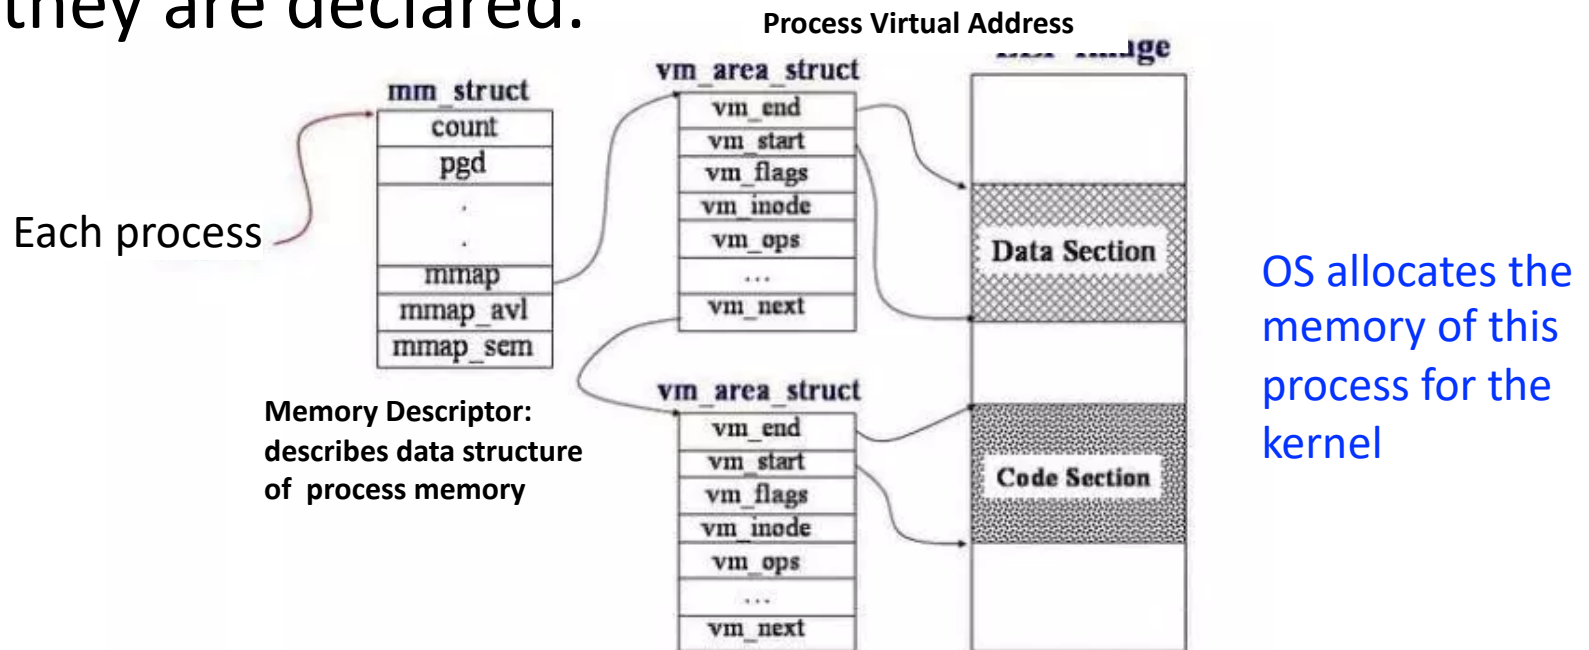
- Various memory models exist for multitasking environments
- Each process requires its own
  - Runtime stack for function calls
  - Storage for local variables
  - Copy of arguments
- A process *may* have private heap storage as well



<http://www.lynx.com/posix-processes-name-spaces-and-virtual-memory/mmu00a/>

# Consequence for Programmers

- A copy of function arguments and local variables are associated with each process executing a particular function, *not* with the code in which they are declared.

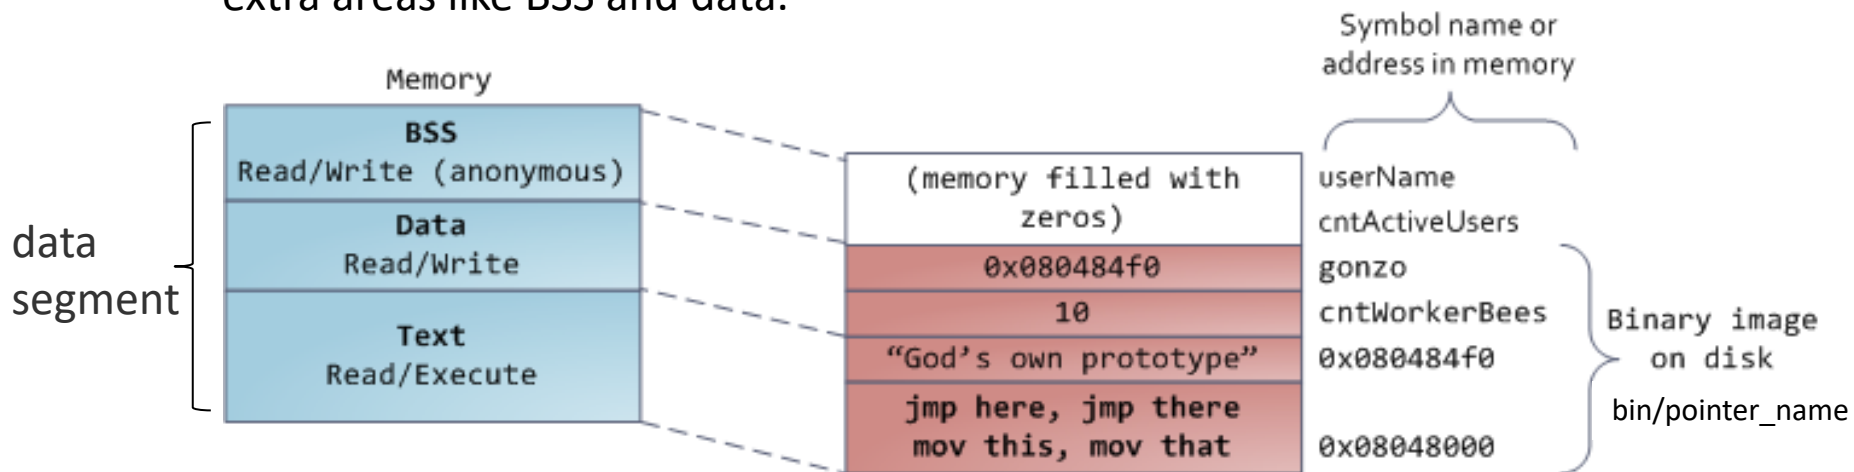


Linux: Management of Process Memory



# Example: How Kernel Manages Memory (1/3)

- The contents of a *pointer* - a 4-byte memory address - live in the data segment
- The **text** segment: the actual string is read-only and stores all your code
  - Maps your binary file in memory
  - Writes to this area earn your program as “**Segmentation Fault**” (Help prevent pointer bugs)
- This diagram shows the segment and the example variables
  - A segment may contain many areas. For example, each memory mapped file normally has its own area in the mmap segment, and dynamic libraries have extra areas like BSS and data.



# Example: How Kernel Manages Memory (2/3)

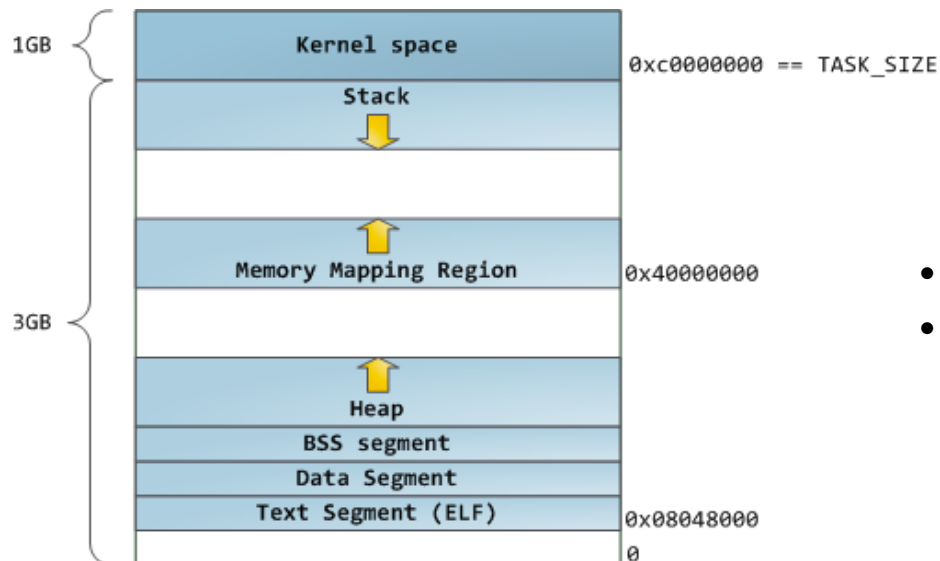
- Examine binary images in Ubuntu using the [nm](#) and [objdump](#) commands to display symbols, their addresses, segments, and so on.

Data structure for process descriptor

```
struct task_struct {  
    ...  
    struct mm_struct  
*mm; /*Memory Descriptor data struct*/  
    ...  
}
```

struct mm\_struct declared

```
struct mm_struct {  
    ...  
    struct vm_area_struct *mmap;  
    /* list of VMAs */  
    ...  
}
```



- The virtual address layout
- A process in Linux is called a task

# Example: How Kernel Manages Memory (3/3)

- Within memory descriptor for managing program memory: the set of virtual memory areas and the page tables

- Data structure of Virtual Address Space (VMA)

```
struct vm_area_struct{
    struct mm_struct * vm_mm;

    /* record VMA area start address*/
    unsigned long vm_start;
    /* record VMA area end address*/
    unsigned long vm_end;
    /* points to the next VMA area data structure*/
    struct vm_area_struct *vm_next;
    ...
}
```

- Goal: VMA is more efficient to manage memory (paging)
- Memory mapping from Virtual to Physical memory: implement by using red-black trees instead of linked list

# **AN OPERATING SYSTEM FROM THE INSIDE**


# Operating System

- Well-understood subsystems
- Many subsystems employ heuristic policies
  - Policies can conflict
  - Heuristics can have corner cases
- Complexity arises from interactions among subsystems
- Side-effects can be
  - Unintended
  - Unanticipated

# Building an Operating System

- The intellectual challenge comes from the “system,” not from individual pieces
- Structured design is needed
- It can be difficult to understand the consequences of choices
- We will use a hierarchical *microkernel* design to help control complexity

1. The kernel is broken down into separate processes
2. Run in kernel/user
3. All servers in different address spaces
4. Servers invoke "services" by sending messages



1. Monolithic kernel is a single large process
2. Single static binary file
3. All services in kernel space
4. Invoke functions directly

# Xinu: Major OS Components

- **Process manager:** decide how to allocate CPU, keep track of status of each process
- **Memory manager:** in charge of main memory, such as checking the validity of each request for memory space
- **Device manger:** monitor every device, channel, and control unit
- Clock (time) manager
- **File manager:** track every file, including data files, assembler, compilers and application programs
- **Interprocess communication:** allow process to communicate with each other
- **Intermachine communication**
- **Assessment and accounting:** user/system, statistics


# Xinu: Multilevel Structure

- Organizes components
- Controls interactions among subsystems
- Allows a system to be understood and built incrementally
- Differs from traditional layered approach
- Will be employed as the design paradigm throughout the text and course



# Multilevel vs. Multilayered Organization (1/2)

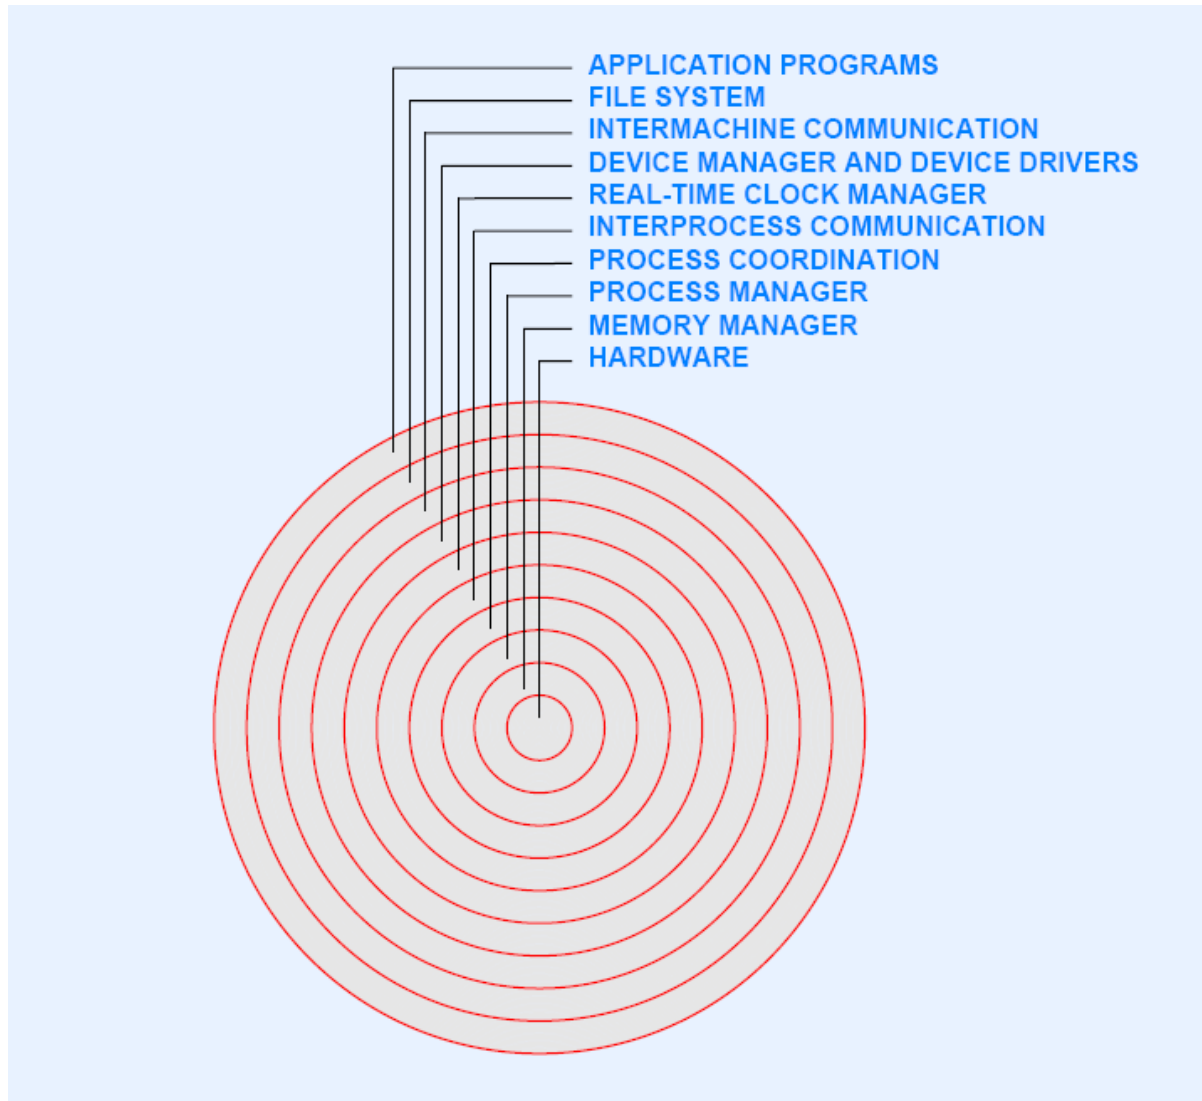
- Multilayer structure
  - Visible to user as well as designer
  - Each layer uses layer directly beneath
  - Involves protection as well as data abstraction
  - Examples
    - Internet protocol layering
    - MULTICS layered security structure
  - Can be inefficient



# Multilevel vs. Multilayered Organization (2/2)

- Multilevel structure
  - Form of data abstraction
  - Used during system construction
  - Helps designer focus attention on one aspect at a time
  - Keeps policy decisions independent
  - Allows given level to use *all* lower levels
  - Efficient

# Multilevel Structure of Xinu



# How to Build an OS

- Work one level at a time
- Identify a service to be provided
- Begin with a *philosophy*
- Establish *policies* that follow the philosophy
- Design *mechanisms* that enforce the policies
- Construct an *implementation* for specific hardware

# Design Example

- Example: access to I/O
- Philosophy: “fairness”
- Policy: FCFS resource access
- Mechanism: queue of requests (FIFO)
- Implementation: program written in C

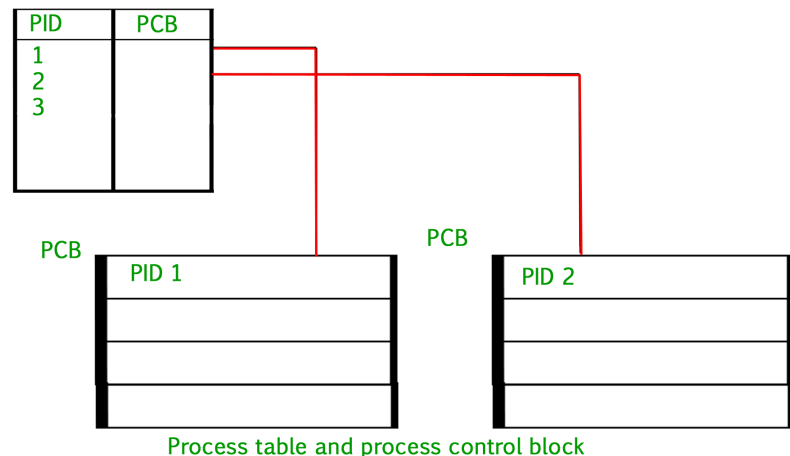
# **DATA STRUCTURE IN XINU: LIST MANIPULATION**

# Queues and Lists

- Fundamental throughout an operating system
- Various forms
  - FIFOs
  - Priority lists
  - Ascending and descending order
  - Event lists ordered by time of occurrence
- Operations
  - *Insert* item
  - *Extract* “next” item
  - *Delete* arbitrary item

# Lists and Queues in Xinu

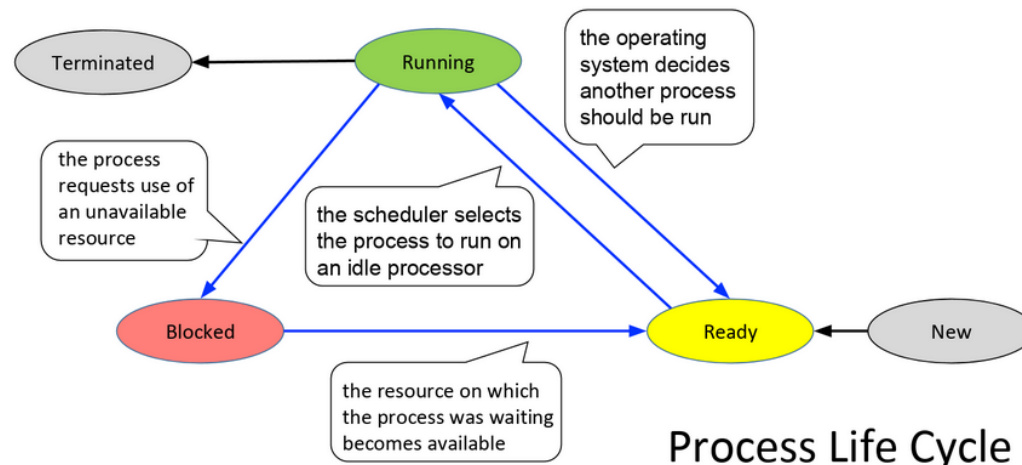
- Important ideas
  - Many lists store processes
  - A process is known by an integer *process ID*
  - A list stores a set of process IDs
- A single data structure can be used to store many types of lists





# Linked Lists in the OS

- Manipulating lists of processes is an important operation in the OS
  - A process's lifecycle consists of moving between, and in, queues and lists
- Xinu implements a unified approach to list management
  - All list management uses this common infrastructure
  - Common functions to create a new list, insert an element at the end of the list, insert or remove from the middle, remove an item from the front

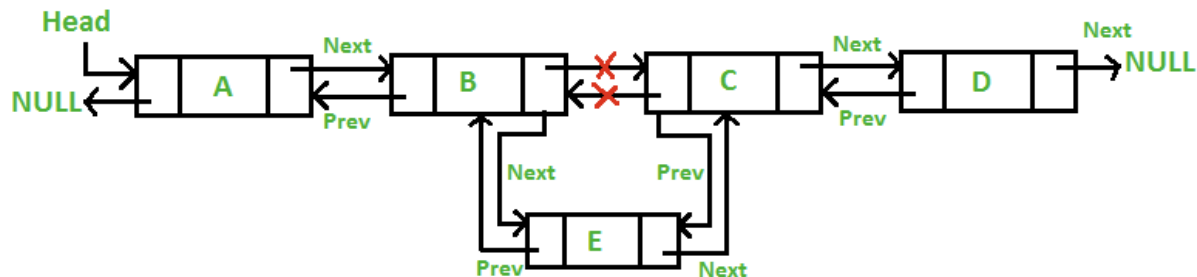


# Unified List Storage in Xinu

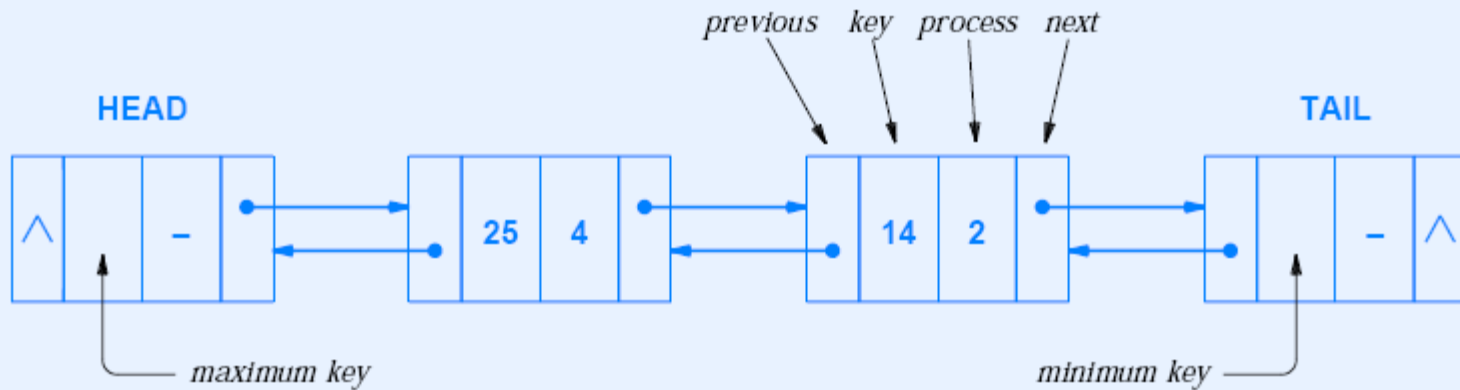
- The process manager handles processes
  - A process moves among the lists frequently
  - At any time, a process is only in one list
- Rather than store all the information about a process, the process manager is free to store only the *process ID (PID)* or *Thread ID (TID)* in a list
  - So when we refer to putting a process on a list, it really means *putting the PID* there – the process control block need not move
- Unified implementation means that not every subsystem uses all the list features

# List Properties

- All lists are doubly-linked – each node points to its predecessor and successor
- Each node stores a *key* as well as a *process ID*, even though a key is not used in a FIFO list
- Each list has a head and tail; the head and tail nodes have the same shape as other nodes
- Non-FIFO lists are ordered in descending order
- The key value in a head node is the *maximum* integer used as a key, and the key value in the tail node is the *minimum* integer used as a key

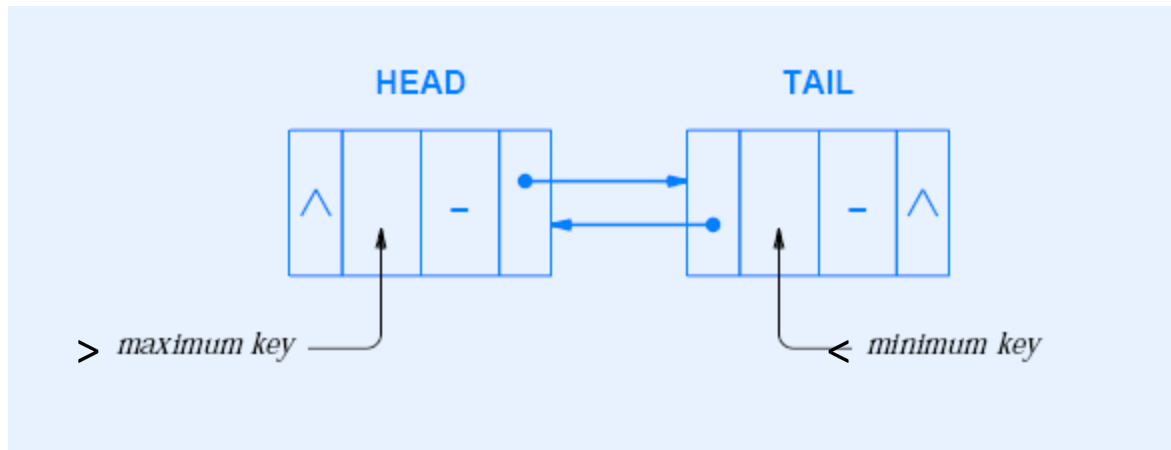


# Conceptual List Structure



- Example list contains two processes, 2 and 4
- Process 4 has key 25
- Process 2 has key 14

# Pointers in an Empty List




- Head and tail linked
- Eliminates special cases for insertion or deletion

# Reducing List Size

- Pointers can mean a large memory footprint
- Important concept: Compact memory usage
  - A process can appear on at most one list at any time
- Techniques used to reduce the size of Xinu lists
  - Relative pointers
  - Implicit data structure
- Most OS place a fixed upper bound on the number of processes
  - In Xinu, constant NPROC specifies the maximum number of processes each user can create, and process identifiers range from 0 through NPROC – 1.

# Relative Pointers

- Store list elements in an array (contiguous memory locations)
  - Each item in array is one node
  - Use *array index* instead of address to identify a node
- Relative pointers
  - Given that there is some (small) fixed number of processes (NPROC)
  - One might use a pointer in this situation, which is 4 bytes (32-bit architecture)
    - Ex: NPROCS < 62, we only need 6 bits
  - Allocate the nodes in a contiguous array and use the array index as a “pointer”

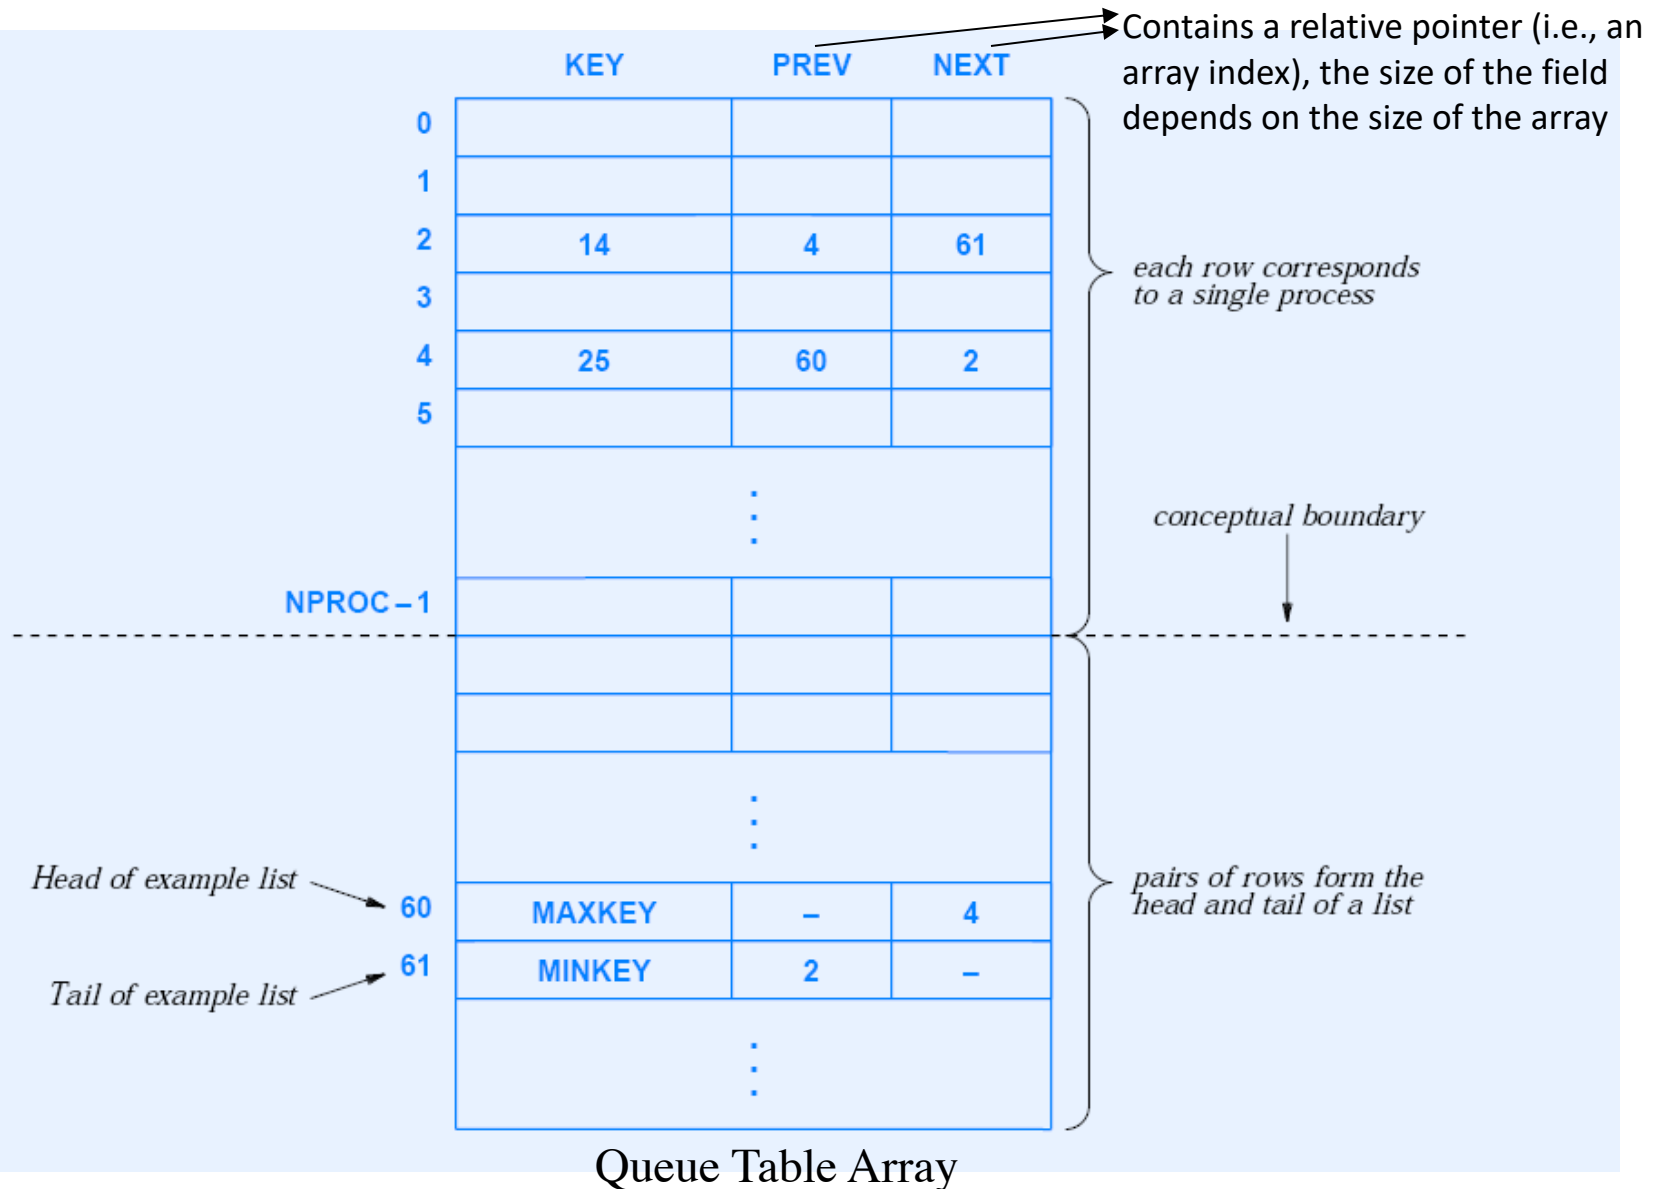


# Implicit data structure

- Omit the process ID field from all nodes
- Because a process can only be in one list, we can use the list position to indicate the ID
- To omit the PID, use an array and use the  $i^{\text{th}}$  element of the array for process  $i$
- Implicit data structure
  - Let  $NPROC$  be the number of processes in the system
  - Assign process IDs 0 through  $NPROC - 1$
  - Let  $i^{\text{th}}$  element of array correspond to process  $i$ , for  $0 \leq i < NPROC$
  - Store heads and tails in same array at positions  $NPROC$  and higher



# Illustration of Xinu List Structure



# Implementation

- Queue data structure
  - Consists of a single array named *queuetab*
  - Is global and available throughout entire OS
- Functions used to manipulate queues
  - Include tests, such as *isempty*, as well as insertion and deletion operations
  - Implemented with inline functions when possible
- Example code shown after discussion of types

# A Question About Types in C

- K&R [1] C defines *short*, *int*, and *long* to be machine-dependent
- ANSI C leaves *int* as a machine-dependent type
- A programmer can define type names
- Question: should a type specify
  - The purpose of an item?
  - The size of an item?
- Example: should a process ID type be named
  - *processid\_t* to indicate the purpose?
  - *int32* to indicate the size?

# Type Names Used in Xinu

- Xinu uses a compromise to encompass both *purpose* and *size*
- Example: consider a variable that holds an index into *queuetab*
- The type name can specify
  - That the variable is a queue table index
  - That the variable is a 16-bit signed integer
- Xinu uses the type name *qid16* to specify both

# Definitions from queue.h (part 1)

```
/* queue.h - firstid, firstkey, isempty, lastkey, nonempty */

/* Queue structure declarations, constants, and inline functions */

/* Default # of queue entries: 1 per process plus 2 for ready list plus */
/*                               2 for sleep list plus 2 per semaphore */
#ifndef NQENT
#define NQENT (NPROC + 4 + NSEM + NSEM)
#endif

#define EMPTY      (-1)          /* Null value for qnext or qprev index */
#define MAXKEY     0x7FFFFFFF    /* Max key that can be stored in queue */
#define MINKEY     0x80000000    /* Min key that can be stored in queue */

struct gentry {                  /* One per process plus two per list */
    int32 qkey;                  /* Key on which the queue is ordered */
    qid16 qnext;                 /* Index of next process or tail */
    qid16 qprev;                /* Index of previous process or head */
};

extern struct gentry queuestab[];
```

← allocates enough space for each process

# Definitions from queue.h (part 2)

```
/* Inline queue manipulation functions */
```

```
#define queuehead(q)      (q)
#define queueetail(q)    ((q) + 1)
#define firstid(q)       (queuetab[queuehead(q)].qnext)
#define lastid(q)        (queuetab[queueetail(q)].qprev)
#define isempty(q)       (firstid(q) >= NPROC)
#define nonempty(q)      (firstid(q) < NPROC)
#define firstkey(q)      (queuetab[firstid(q)].qkey)
#define lastkey(q)       (queuetab[ lastid(q)].qkey)
```

```
/* Inline to check queue id assumes interrupts are disabled */
```

```
#define isbadqid(x)       (((int32)(x) < 0) || (int32)(x) >= NQENT-1)
```


# List manipulation functions

- Look at the list manipulation functions
  - *enqueue* and *dequeue* (in *queue.c*)
  - *getfirst*, *getlast*, and *getitem* (in *getitem.c*)
  - *insert* (in *insert.c*)
  - *newqueue* (in *newqueue.c*)

# Code for Insertion and Deletion from a Queue (part 1)

```
/* queue.c - enqueue, dequeue */
#include <xinu.h>
struct gentry queuetab[NQENT];          /* Table of process queues */
/*-----
 * enqueue - Insert a process at the tail of a queue
 *-----
 */
pid32 enqueue(
    pid32 pid,                          /* ID of process to insert */
    qid16 q                             /* ID of queue to use */
)
{
    int tail, prev;                    /* Tail & previous node indexes */
    if (isbadqid(q) || isbadpid(pid)) {
        return SYSERR;
    }
    tail = queuetail(q);
    prev = queuetab[tail].qprev;
    queuetab[pid].qnext = tail;
    queuetab[pid].qprev = prev;
    queuetab[prev].qnext = pid;
    queuetab[tail].qprev = pid;
    return pid;
}
```


/\* Insert just before tail node \*/





# Code for Insertion and Deletion from a Queue (part 2)

```
/*-----  
 * dequeue - Remove and return the first process on a list  
 *-----  
 */  
pid32 dequeue(  
    qid16 q                                /* ID queue to use */  
)  
{  
    pid32 pid;                            /* ID of process removed */  
  
    if (isbadqid(q)) {  
        return SYSERR;  
    } else if (isempty(q)) {  
        return EMPTY;  
    }  
    pid = getfirst(q);  
    queuetab[pid].qprev = EMPTY;  
    queuetab[pid].qnext = EMPTY;  
    return pid;  
}
```



# Code for Insertion in an Ordered List (part 1)

```
/* insert.c - insert */
#include <xinu.h>
/*-----
 * insert - Insert a process into a queue in descending key order
 *-----
 */
status insert(
    pid32 pid,                /* ID of process to insert */
    qid16 q,                  /* ID of queue to use */
    int32 key                  /* Key for the inserted process */
)
{
    int16 curr;               /* Runs through items in a queue*/
    int16 prev;               /* Holds previous node index */

    if (isbadqid(q) || isbadpid(pid)) {
        return SYSERR;
    }
    curr = firstid(q);
    while (queuetab[curr].qkey >= key) {
        curr = queuetab[curr].qnext;
    }
}
```

# Code for Insertion in an Ordered List (part 2)

```
/* Insert process between curr node and previous node */
```

```
prev = queue[tab[curr]].qprev;    /* Get index of previous node */  
queue[tab[pid]].qnext = curr;  
queue[tab[pid]].qprev = prev;  
queue[tab[pid]].qkey = key;  
queue[tab[prev]].qnext = pid;  
queue[tab[curr]].qprev = pid;  
return OK;
```

```
}
```

# Accessing an Item in a List (part 1)

```
/* getitem.c - getfirst, getlast, getitem */
```

```
#include <xinu.h>
```

```
/*-----  
 * getfirst - Remove a process from the front of a queue  
 *-----  
 */  
pid32 getfirst(  
    qid16 q                /* ID of queue from which to */  
)                        /* Remove a process (assumed */  
                          /* valid with no check) */  
{  
    pid32 head;  
  
    if (isempty(q)) {  
        return EMPTY;  
    }  
  
    head = queuehead(q);  
    return getitem(queuestab[head].qnext);  
}
```

# Accessing an Item in a List (part 2)

```
/*-----  
 * getlast - Remove a process from end of queue  
 *-----  
 */  
pid32 getlast(  
    qid16 q          /* ID of queue from which to */  
)                  /* Remove a process (assumed */  
                    /* valid with no check) */  
{  
    pid32 tail;  
  
    if (isempty(q)) {  
        return EMPTY;  
    }  
  
    tail = queueutil(q);  
    return getitem(queueutil[tail].qprev);  
}
```

# Accessing an Item in a List (part 3)

```
/*-----  
 * getitem - Remove a process from an arbitrary point in a queue  
 *-----  
 */  
pid32 getitem(  
    pid32 pid                                /* ID of process to remove */  
)  
{  
    pid32 prev, next;  
  
    next = queuetab[pid].qnext;              /* Following node in list */  
    prev = queuetab[pid].qprev;              /* Previous node in list */  
    queuetab[prev].qnext = next;  
    queuetab[next].qprev = prev;  
    return pid;  
}
```

# Allocating a New List

```
/* excerpt from newqueue.c */
```

```
qid16 newqueue(void)
{
    static qid16    nextqid=NPROC; /* Next list in queuetab to use */
    qid16          q;              /* ID of allocated queue      */

    q = nextqid;
    if (q > NQENT) {                /* Check for table overflow */
        return SYSERR;
    }

    nextqid += 2;                   /* Increment index for next call*/

    /* Initialize head and tail nodes to form an empty queue */

    queuetab[queuehead(q)].qnext = queuetail(q);
    queuetab[queuehead(q)].qprev = EMPTY;
    queuetab[queuehead(q)].qkey  = MAXKEY;
    queuetab[queuetail(q)].qnext = EMPTY;
    queuetab[queuetail(q)].qprev = queuehead(q);
    queuetab[queuetail(q)].qkey  = MINKEY;
    return q;
}
```

# Summary

- Operating system supplies set of services
- System calls provide interface between OS and application
- Concurrency is fundamental concept
  - Between I /O devices and processor
  - Between multiple computations
- Process is OS abstraction for concurrency
- Process differs from program or function
- You will learn how to design and implement system software that supports concurrent processing



# Summary

(continued)

- OS has well-understood internal components
- Complexity arises from interactions among components
- Multilevel approach helps organize system structure
- Design involves inventing policies and mechanisms that enforce overall goals
- Xinu includes a compact list structure that uses relative pointers and an implicit data structure to reduce size
- Xinu type names specify both purpose and data size

# References

1. Brian Kernighan and Dennis Ritchie published the first edition of *The C Programming Language* in 1978, known to C programmers as "K&R", as an informal specification of the language.
2. Introduction to Programming Systems: [Systems Calls and Standard](#) by Professor Jennifer Rexford from Princeton University.