


# CSCI 8530

# Advanced Operating Systems

Part 6

Inter-Process Communication


# Location of Inter-Process Communication in the Hierarchy




# Inter-process Communication

- A process can be of two type: Independent process (not affected by the execution of other processes ), Cooperating process.
- A mechanism allows processes to communicate each other and synchronize their action
- Used for
  - Exchange of (nonshared) data
  - Process coordination
- Use these two ways: *Shared Memory, Message passing*

- Communication takes place by way of messages exchanged among the cooperating processes.




(a) Message Passing



(b) Shared Memory

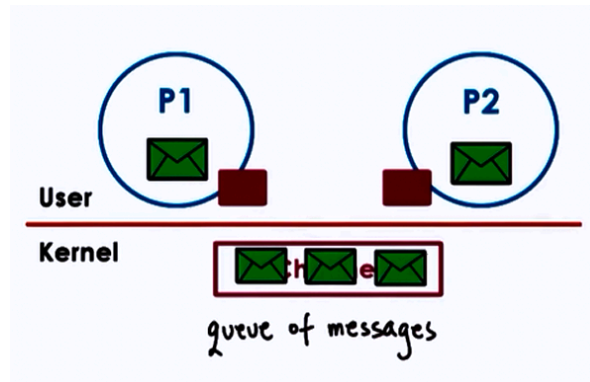
- A region of memory which is shared by cooperating processes gets established.
- Processes can exchange information by reading and writing all the data to the shared region.

# Forms of Message Passing IPC




## Pipes

Carry byte stream  
between 2 process



## Message queues

Carry "messages" among  
processes



## Sockets

- `send()` and `recv()` : pass message buffers
- `socket()` : create kernel level socket buffer


# Two Approaches to Message Passing

- Approach #1
  - Message passing is one of many services
  - Messages are separate from I /O and process synchronization services
  - Messages implemented using lower-level mechanisms, such as semaphores
- Approach #2
  - The entire operating system is *message-based*
  - Messages, not function calls, provide the fundamental building block
  - Messages, not semaphores, used for process synchronization

# Typical Message-Passing Functions

The following functions and data items are characteristic in a typical message-passing environment:

- *Source* and *Destination* addresses must identify the machine and the process being addressed.
- *message* is a generic data item to be delivered.
- Use two system call
  - send (destination, &message): sends message to destination without blocking; the sender's address is usually included
  - receive (source, &message): receives the next sequential message from source, blocking until it is available.



# Design of a Message Passing Facility

- To understand the issues, we will begin with a trivial message passing facility
- We want to allow a process to send a message directly to another process
- In principle,  
the design should be straightforward
- In practice,  
many design decisions arise

# Message Passing Design Decisions

- Are messages fixed or variable size?
- What is the maximum message size?
- How many messages can be outstanding at a given time?
- Where are messages stored?
- How is a recipient specified?
- Does a receiver know the sender's identity?
- Are replies supported?
- Is the interface synchronous or asynchronous?



# Synchronous vs. Asynchronous Interface (1/3)

Can message passing also replace synchronization primitives such as semaphores?

- The answer depends on the implementation of message passing

# Synchronous vs. Asynchronous Interface (2/3)

1. Synchronous interface: To receive a message in a synchronous system, a process calls a system function, and the call does not return until a message arrives
  - Blocks until the operation is performed
  - Easy to understand / program
  - Extra processes can be used to obtain asynchrony
  - Using message passing to implement mutual exclusion is possible

# Synchronous vs. Asynchronous Interface (3/3)

2. Asynchronous interface: an asynchronous message passing system either requires a process to *poll* or a mechanism that allows OS to stop a process temporarily
  - Process starts an operation
  - Initiating process continues execution
  - Notification
    - Arrives when operation completes
    - May entail abnormal control flow (e.g., software interrupt or “callback” mechanism)
  - Polling can be used to determine status
  - Additional overhead or complexity, but convenient

# Two Forms of Message Passing in Xinu

- A completely synchronous paradigm and a partially asynchronous paradigm
- Provide *direct* and *indirect* message delivery:
  - One provides a direct exchange of messages among processes (Ch. 8)
  - The other arranges for messages to be exchanged through rendezvous point (Ch11)

# Why is a Message Passing Facility So Difficult to Design?

- Interacts with
  - Process coordination subsystem
  - Memory management subsystem
- Affects user's perception of system

# Message passing system in Xinu

- The message passing facility follows three guidelines:
  - *Limited message size*. The system limits each message to a small, fixed size.
  - *No message queues*. The system permits a given process to store only one unreceived message per process at any time.
  - *First message semantics*. If several messages are sent to a given process before the process receives any of them, only the first message is stored and delivered; subsequent senders do not block.
    - First message semantics: for determining which of several events completes first.

# Example Inter-process Message Passing Design (1/3)

- Simple, low-level mechanism
- Direct process-to-process communication
- One-word messages
- Message stored with receiver
- One-message buffer
- Synchronous, buffered reception
- Asynchronous transmission and “reset” operation

# Example Inter-process Message Passing Design (2/3)

- Three system calls manipulate messages
  - send*(msg, pid);
  - msg = *receive*();
  - msg = *recvclr*();
- *Send* transmits message to specified process
- *Receive* blocks until a message arrives
- *Recvclr*: Remove existing message, if one has arrived, but does not block
  - Receive a message → return the message (like *receive*)
  - No message is waiting → return OK
- Message stored in *receiver's* process table entry



# Example Inter-process Message Passing Design (3/3)

- First-message semantics
  - First message sent to a process is stored until it has been received
  - Subsequent attempts to send fail
- Idiom

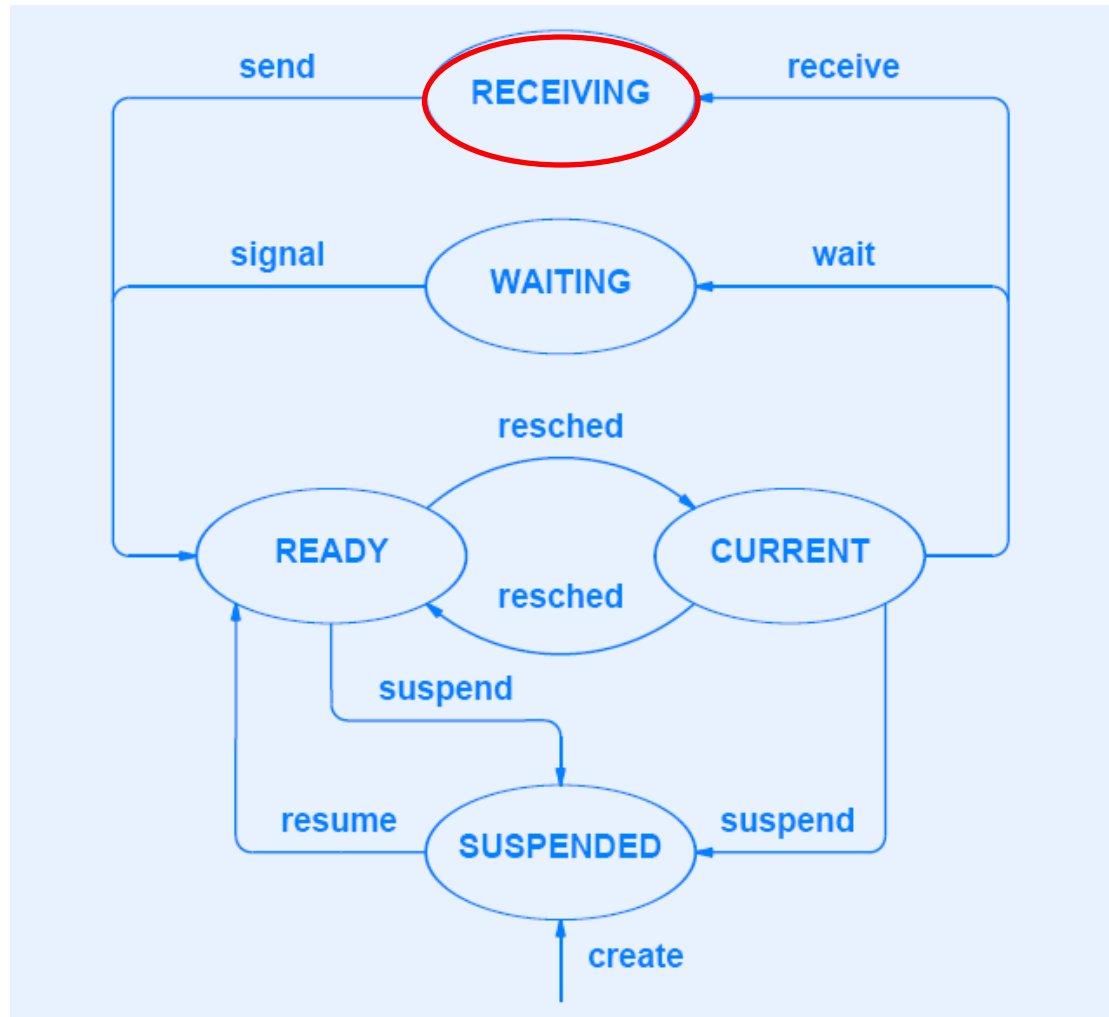
```
recvclr();    /* prepare to receive a message */  
... /* allow other processes to send messages */  
msg = receive();
```
- Above code returns first message that was sent, even if a high priority process sends later

# Process State for Message Reception

In what state should a process be while waiting for a message?

- While receiving a message, a process is NOT
  - Executing
  - Ready
  - Suspended
  - Waiting on a semaphore
- Therefore, a new state is needed for message passing
- Named *RECEIVING*, with the symbolic constant `PR_RECV`
- Entered when *receive* called

# State Transitions with Message Passing



The symbolic  
constant PR\_RECV

# Xinu Code for Message Transmission (1/2)

```
/* send.c - send */
```

A message cannot be stored **in the sender's memory**

```
#include <xinu.h>
```

```
/*-----  
 * send - Pass a message to a process and start recipient if waiting  
 *-----  
 */  
syscall send(  
    pid32 pid, /* ID of recipient process */  
    umsg32 msg /* Contents of message */  
)  
{  
    intmask mask; /* Saved interrupt mask */  
    struct procent *prptr; /* Ptr to process' table entry */  
  
    mask = disable();  
    if (isbadpid(pid)) {  
        restore(mask);  
        return SYSERR;  
    }  
    prptr = &proctab[pid];  
    if ((prptr->prstate == PR_FREE) || prptr->prhasmsg) {  
        restore(mask);  
        return SYSERR;  
    }  
}
```

Check to ensure the recipient does not have a message outstanding

# Xinu Code for Message Transmission

## (2/2)

Restrictions on the size of messages: The implementation reserve space for one message

```
prptr->prmsg = msg;  
prptr->prhasmsg = TRUE;
```

```
/* Deliver message */  
/* Indicate message is waiting */
```

```
/* If recipient waiting or in timed-wait make it ready */
```

```
if (prptr->prstate == PR_RECV) { /* If the recipient is waiting */  
                                /* for the arrival of a message */
```

```
    ready(pid);
```

```
} else if (prptr->prstate == PR_RECTIM) {
```

```
    unsleep(pid);
```

```
    ready(pid);
```

Remove the process from the queue of sleeping processes

```
}
```

```
restore(mask);
```

```
/* Restore interrupts */
```

```
return OK;
```

```
}
```

- Note: we will discuss PR\_RECTIM (receive-with-timeout) later

# Xinu Code for Message Reception

```
/* receive.c - receive */
```

```
#include <xinu.h>
```

```
/*-----  
 * receive - wait for an incoming message and return the message to the caller  
 *-----  
 */
```

```
umsg32 receive(void)
```

```
{
```

```
    intmask mask;                /* Saved interrupt mask */  
    struct procent *prptr;        /* Ptr to process' table entry */  
    umsg32 msg;                  /* Message to return */
```

```
    mask = disable();
```

```
    prptr = &proctab[currpid];
```

Determine whether a message is waiting

```
    if (prptr->prhasmsg == FALSE) { /* If no message has arrived */
```

```
        prptr->prstate = PR_RECV;
```

```
        resched(); /* Block until message arrives */
```

```
    }
```

```
    msg = prptr->prmsg; /* Retrieve message */
```

```
    prptr->prhasmsg = FALSE; /* Reset message flag */
```

```
    restore(mask);
```

```
    return msg;
```

```
}
```

# Xinu Code for Clearing Messages

```
/* recvclr.c - recvclr */

#include <xinu.h>

/*-----
 * recvclr - Clear incoming message, and return message if one waiting
 *-----
 */
umsg32 recvclr(void)
{
    intmask mask;                /* Saved interrupt mask */
    struct procent *prptr;        /* Ptr to process' table entry */
    umsg32 msg;                   /* Message to return */

    mask = disable();
    prptr = &proctab[currpid];
    if (prptr->prhasmsg == TRUE) {
        msg = prptr->prmsg;        /* Retrieve message */
        prptr->prhasmsg = FALSE;    /* Reset message flag */
    } else {
        msg = OK;
    }
    restore(mask);
    return msg;
}
```

# Summary

- Inter-process communication
  - Implemented by message passing
    - A low-level mechanism provides direct communication among processes
    - A high-level mechanism uses rendezvous points.
  - Can be synchronous or asynchronous
- Low-level mechanism in Xinu:
  - Limits the message size to a single word
  - Restricts each process to at most one outstanding message
  - Use first-message semantics
- Synchronous interface is the simplest
- Xinu uses synchronous reception and asynchronous transmission