


CSCI 8530

Advanced Operating Systems

Part 10

Clock and Timer Management


Location of Device Management in the Hierarchy



Clock Hardware (1/3)

1. Processor clock

- Controls processor rate
- Often cited as processor speed
- Usually not visible to the OS



2. Time-of-day clock (chronometer)

- Can be set or read by the processor
- Operates independently from processor

Clock Hardware (2/3)

3. Real-time clock

- Pulses regularly
- Called *programmable* if rate can be controlled by OS
- Interrupts processor on each pulse
- Does not count pulses; interrupt is lost if missed
 - Means OS cannot leave interrupts disabled for more than one clock tick or real-time clock will be inaccurate (interrupts/pulses will be missed)

Note: some real-time clock hardware has *counter* for missed interrupts, which allows OS to correct counts later

Clock Hardware (3/3)

4. Interval timer (count-down timer)

- Hardware device that operates asynchronously
- An internal real-time clock and a counter
- Timeout value T is set by processor, and device interrupts T time units later
- Processor may be able to find time remaining
- *Programmable* needs two registers
 - The “period” register: how much time should pass between consecutive interrupts
 - Generate multiple interrupts, copied to “remaining”
 - The “remaining” register: how much time remains before the next interrupt
 - Decrement at a specified rate

TIME MANAGEMENT

Timed Events

- *Hard real-time* system requires each event to occur at (or no later than) an exact time
- *Soft real-time* system requires an event to occur on or *occasionally* after the specified time
- In theory, an OS may need to handle many event types
- In practice, only two basic event types suffice

Two Principle Types of Timed Events

- *Preemption event*
 - Implements *timeslicing* by switching the processor among ready processes
 - Guarantees a given process cannot run forever
 - A preemption event is scheduled during a context switch
 - Cancellation is important (few processes exhaust their timeslice)
- *Sleep event*
 - Move to a new state
 - Scheduled by a process
 - Delays the calling process a specified time

Time-slicing Tradeoff

- How large should a timeslice be?
- Small granularity
 - Advantage: guarantees fairness because processes proceed at approximately equal rate
 - Disadvantage: increased rescheduling overhead
- Large granularity
 - Advantage: lower rescheduling overhead
 - Disadvantage: unfair because one process may be far ahead of another

Timeslicing and Conventional Applications

Most applications are I/O bound, which means the application is likely to perform an operation that takes the process out of the current state before its timeslice expires.

Without a preemptive capability, an operating system cannot regain control from a process that executes an infinite loop.

Managing Real-time Clock Events


- Must be efficient
 - Clock interrupts occur frequently and continually
 - Set of timed events examined on each clock interrupt
 - Should avoid searching a list (avoid $O(N)$ complexity)
- Mechanism
 - Keep timed events on a linked list
 - One item appears on the list for each outstanding event
 - Called *event queue* (aka *callout queue*)

Delta List

- Data structure used for timed events
- Items on the list are ordered by time of occurrence
- For efficiency, list stores *relative* times
- The key in an item stores the difference (*delta*) between the time for the event and time for the previous event
- The key in first event stores the delta from “now”
- In UNIX, items are called “callouts”

Delta List Example

- Assume events for processes *A* through *D* will occur 6, 12, 27, and 50 ticks from now
- The delta keys are 6, 6, 15, and 23



Real-time Clock Processing in Xinu

- Sleep queue
 - Delta list of delayed processes
 - Each node on the list corresponds to a sleeping process
- Global variable *sleepq* contains the ID of the sleep queue
- Clock interrupt handler
 - Decrements preemption counter
 - Reschedules if timeslice has expired
 - Processes event list

Keys on the Xinu Sleep Queue

- Processes on *sleepq* are ordered by time at which they will awaken
- Each key tells the number of clock ticks that the process must delay beyond the preceding one on the list
- Relationship must be maintained whenever an item is inserted or deleted

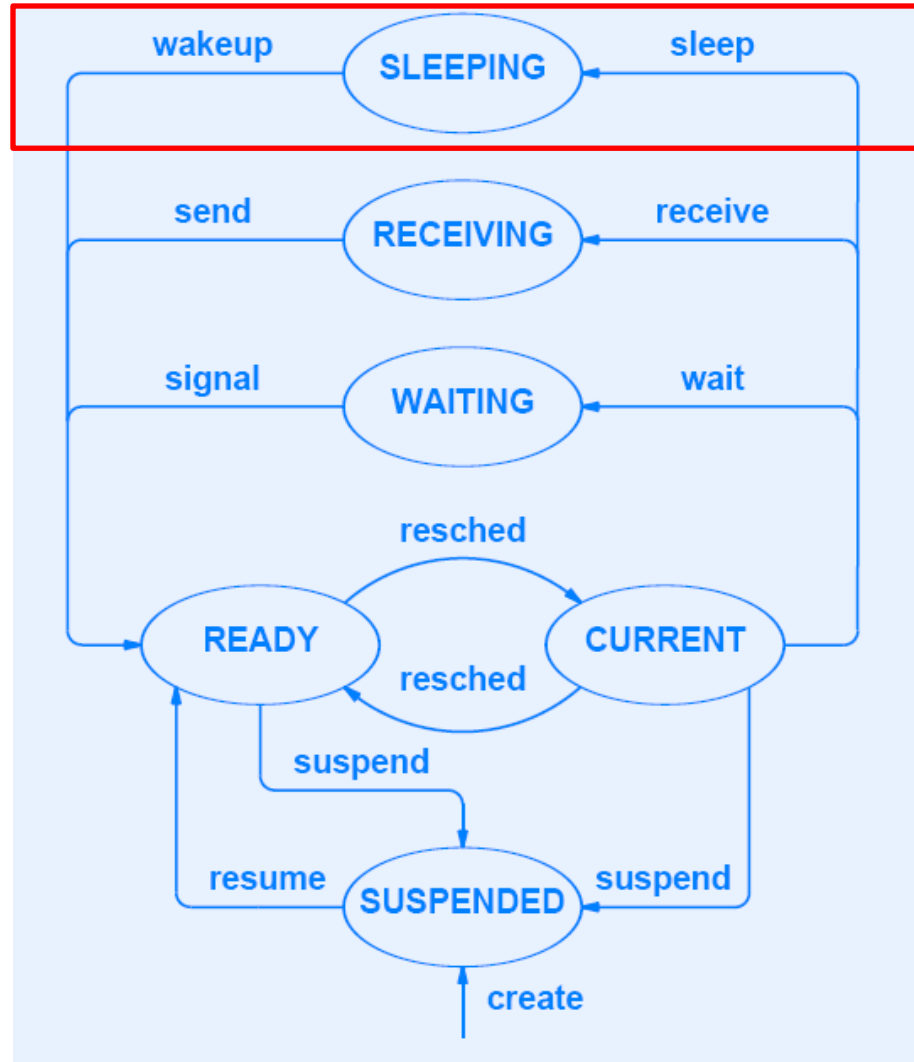
Real-time Delay and Clock Resolution

- Process calls *sleep* to delay
- Question: what resolution should be used for sleep?
 - Humans typically think in seconds or minutes
 - Real-time applications may need millisecond accuracy or better (if available)

Xinu Sleep Primitives and Resolution

- Set of primitives accommodates range of possible resolutions
 - sleep – delay given in seconds
 - sleep10 – delay given in tenths of seconds
 - sleep100 – delay given in hundredths of seconds
 - sleepms – delay given in milliseconds
- Note: some hardware does not support all resolutions

New Process State for Sleeping Processes



sleepms inserts the process into the delta list of sleeping processes

Xinu Sleep Function (part 1)

```
/* sleep.c - sleep sleeps */
```

```
#include <xinu.h>
```

```
#define MAXSECONDS 4294967 /* Max seconds per 32-bit msec */
```

```
/*-----  
 * sleep - Delay the calling process n seconds  
 *-----  
 */  
syscall sleep(  
    uint32 delay /* Time to delay in seconds */  
)  
{  
    if (delay > MAXSECONDS) {  
        return SYSERR;  
    }  
    sleeps(1000*delay);  
    return OK;  
}
```

Xinu Sleep Function (part 2)

```
/*-----  
 * sleepms - Delay the calling process n milliseconds  
 *-----  
 */  
syscall sleepms(  
    uint32 delay /* Time to delay in msec. */  
)  
{  
    intmask mask; /* saved interrupt mask */  
  
    mask = disable();  
    if (delay == 0) {  
        yield();  
        restore(mask);  
        return OK;  
    }  
    /* Delay calling process */  
    if (insertd(currpid, sleepq, delay) == SYSERR) {  
        restore(mask);  
        return SYSERR;  
    }  
    proctab[currpid].prstate = PR_SLEEP;  
    resched();  
    restore(mask);  
    return OK;  
}
```

sleepms uses *insertd* to insert
the current process in the delta
list of sleeping processes

Given by argument key

Inserting an Item on Sleepq

- Current process calls *sleepms* or *sleep* to request delay
- *Sleepms*
 - Underlying function that takes action
 - Inserts current process on *sleepq*
 - Calls *resched* to allow other processes to execute
- Method
 - Walk through *sleepq* (with interrupts disabled)
 - Find place to insert the process
 - Adjust remaining keys as necessary

Xinu Insertd (part 1)

```
/* insertd.c - insertd */
```

```
#include <xinu.h>
```

```
/*-----  
 * insertd - Insert a process in delta list using delay as the key  
 *-----  
 */
```

```
status insertd(                                /* Assumes interrupts disabled */  
    pid32 pid,                                /* ID of process to insert */  
    qid16 q,                                  /* ID of queue to use */  
    int32 key                                 /* Delay from "now" (in ms.) */  
)  
{
```

```
    int32 next;                                /* Runs through the delta list */  
    int32 prev;                                /* Follows next through the list*/  
    
```

```
    if (isbadqid(q) || isbadpid(pid)) {  
        return SYSERR;  
    }
```

Xinu Insertd (part 2)

```
prev = queuehead(q);
next = queuetab[queuehead(q)].qnext;
while ((next != queuetail(q)) && (queuetab[next].qkey <= key)) {
    key -= queuetab[next].qkey;
    prev = next;
    next = queuetab[next].qnext;
}
```

Specify a delay relative to the time at which the predecessor of “next” awakens

/* Insert new node between prev and next nodes */

```
queuetab[pid].qnext = next;
queuetab[pid].qprev = prev;
queuetab[pid].qkey = key;
queuetab[prev].qnext = pid;
queuetab[next].qprev = pid;
if (next != queuetail(q)) {
    queuetab[next].qkey -= key;
}
return OK;
```

Insertd must also subtract the extra delay that the new item introduces from the delay of the rest of the list


```
}
```

Invariant During Sleepq Insertion

At any time during the search, both key and `queuetab[next].qkey` specify a delay relative to the time at which the predecessor of the “next” process awakens.

Testing a Clock

- On some systems, clock hardware is optional
- If optional, OS can test for presence of a clock
 - Initialize clock interrupt vector
 - Enable interrupts
 - Loop “long enough”
 - If interrupt occurs, declare clock present
 - Otherwise, declare no clock present and disable *sleep*



Clock Interrupt Handler

- May need to be optimized
- Decrements preemption counter
 - Calls *resched* if counter reaches zero
- Decrements count of the first element on clock queue
 - Calls *wakeup* if counter reaches zero
- Important notes
 - More than one process may awaken at the same time
 - *Wakeup* should awaken all processes that have zero time remaining before allowing any of them to run

Xinu Wakeup

```
/* wakeup.c - wakeup */
```

```
#include <xinu.h>
```

```
/*-----  
 * wakeup - Called by clock interrupt handler to awaken processes  
 *-----  
 */  
void wakeup(void)  
{  
    /* Awaken all processes that have no more time to sleep */  
  
    resched_cntl(DEFER_START);  
    while (nonempty(sleepq) && (firstkey(sleepq) <= 0)) {  
        ready(dequeue(sleepq));  
    }  
    resched_cntl(DEFER_STOP);  
    return;  
}
```

- Rescheduling is deferred until all processes are awakened

Operation Timeout

- Many operating systems components need a “timeout” feature
- Especially useful in building communication protocols
- Possible approaches
 - Add timeout feature to each system call (difficult)
 - Provide a single facility for timeout
- Xinu uses latter approach with timed message reception

Timed Message Reception

- Implemented with *recvtime()*
- Argument specifies maximum delay
- A call to *recvtime()* returns
 - If a message arrives before the specified timeout
 - After the specified timeout if no message has arrived
- Value *TIMEOUT* is returned if time expires
- Note: timer is cancelled if message arrives

Xinu Recvtime (part 1)

```
/* recvtime.c - recvtime */

#include <xinu.h>

/*-----
 * recvtime - wait specified time to receive a message and return
 *-----
 */
umsg32 recvtime(
    int32 maxwait                /* Ticks to wait before timeout */
)
{
    intmask mask;                /* Saved interrupt mask */
    struct procent *prptr;       /* entry of current process */
    umsg32 msg;                  /* Message to return */

    if (maxwait < 0) {
        return SYSERR;
    }
    mask = disable();
}
```

Xinu Recvtime (part 2)

```
/* schedule wakeup and place process in timed-receive state */

prptr = &proctab[currpid];
if (prptr->prhasmsg == FALSE) { /* If message waiting, no delay */
    if (insertd(currpid,sleepq,maxwait) == SYSERR) {
        restore(mask);
        return SYSERR;
    }
    prptr->prstate = PR_RECTIM; /* A receive with timeout */
    resched();
}

/* Either message arrived or timer expired */

if (prptr->prhasmsg) {
    msg = prptr->prmsg;          /* Retrieve message */
    prptr->prhasmsg = FALSE; /* Reset message indicator */
} else {
    msg = TIMEOUT;
}
restore(mask);
return msg;
}
```

Example Clock Interrupt Handler (part 1)

```
/* clkhandler.c - clkhandler */

#include <xinu.h>

/*-----
 * clkhandler - high level clock interrupt handler
 *-----
 */
void clkhandler()
{
    static uint32 count1000 = 1000; /* Count to 1000 ms */

    /* Decrement the ms counter, and see if a second has passed */
    if((--count1000) <= 0) {
        /* One second has passed, so increment seconds count */
        clktime++; // provide the date (e.g., it is used by the
                  // xinu shell command date)

        /* Reset the local ms counter for the next second */
        count1000 = 1000; // counts from 1000 down to 0
    }
}
```

Manages sleeping processes and Preemption

Store the time in
seconds since the
system booted

clktime++; // provide the date (e.g., it is used by the
// xinu shell command date)

Example Clock Interrupt Handler (part 2)

```
/* Handle sleeping processes if any exist */
```

```
if(!isempty(sleepq)) {
```

```
    /* Decrement the delay for the first process on the */  
    /* sleep queue, and awaken if the count reaches zero */
```

```
    if((--queuetab[firstid(sleepq)].qkey) <= 0) {
```

```
        wakeup();
```

```
    }
```

Remove all processes from the sleep queue that have a zero delay

```
}
```

```
/* Decrement the preemption counter, and reschedule when the */  
/* remaining time reaches zero */
```

```
if ((--preempt) <= 0) {  
    preempt = QUANTUM;
```

```
    resched();
```

```
}
```

```
}
```

Summary (1/2)

- Computer can contain several types of hardware clocks
 - Processor
 - Time of day
 - Real-time
 - Interval timer
- Real-time clock or interval timer used for
 - Preemption
 - Process delay
- OS may need to convert hardware pulse rate to appropriate tick rate

Summary (2/2)

- Delta list provides elegant data structure to store a set of sleeping processes
- Only the key of the first item on the list needs to be updated on each clock tick
- Multiple processes may awaken at the same time; rescheduling is deferred until all have been made ready
- *Recvtime* allows a process to wait a specified time for a message to arrive