


CSCI 8530

Advanced Operating Systems

Part 11

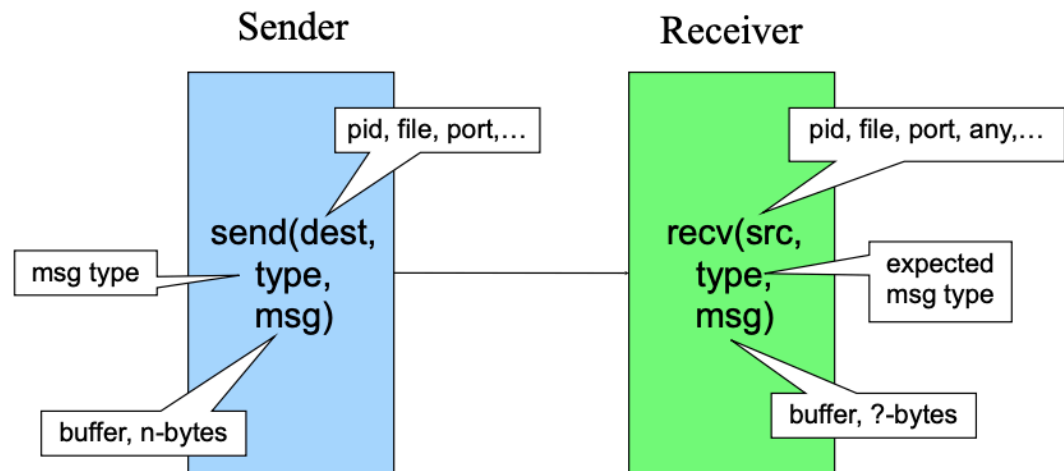
High-level Synchronous Message Passing

Location of Synchronous Message Passing in the Hierarchy



Review of Message Passing Choices


- Potential synchronization
 - Sender blocks
 - Receiver blocks
 - Neither blocks
 - Both block
- Messages outstanding at a given time
 - Arbitrary number
 - Small, fixed number



Review of Message Passing Choices

(continued)

- Use two system call
 - send (destination, &message)
 - receive (source, &message)
- Message storage
 - Associated with sender
 - Associated with receiver
 - Independent of sender and receiver
- Destination
 - A specific process
 - Intermediate pickup point accessible to multiple processes



More Review: Xinu Low-level Message Passing

- Asynchronous (non-blocking) transmission
- Synchronous (blocking) reception
- Asynchronous message clear
- Message buffer holds one message
- Destination is a specific process

Motivations for High-level Message Passing

- Permit synchronous message transfer
 - Block sender until receiver ready
 - Block receiver until sender ready
 - Example: data pipeline
- Make a message available to any process in a set
- Example
 - Concurrent server
 - Set of processes that can handle requests
 - Next process in set handles each incoming request
 - Allows short requests to be serviced quickly

Xinu High-Level Message Passing Mechanism

- Use inter-process communication port (a *port* mechanism) to refer to a rendezvous point
- Separate abstraction, unrelated to low-level message passing
- Part of kernel
- Independent from processes
- Allows arbitrary process to
 - Send messages
 - Receive messages

Port Details

- Port
 - Created dynamically
 - Provides a synchronous interface using a producer and consumer semaphore per port
 - Receiver blocks when port empty
 - Sender blocks when port full
 - Requests are also handled in FIFO order
- At port creation
 - Maximum number of messages is fixed and storage is allocated
 - Semaphores are created
 - Producers and Consumers

Port Declarations

```
/* ports.h - isbadport */
```

```
#define NPORTS    30                /* Maximum number of ports    */
#define PT_MSGS   100               /* Total messages in system   */
#define PT_FREE   1                 /* Port is free                */
#define PT_LIMBO  2                 /* Port is being deleted/reset */
#define PT_ALLOC  3                 /* Port is allocated           */

struct ptnode {                    /* Node on list of messages    */
    uint32 ptmsg;                  /* A one-word message          */
    struct ptnode *ptnext;         /* Pointer to next node on list */
};

struct ptentry {                   /* Entry in the port table     */
    sid32 ptssem;                  /* Sender semaphore            */
    sid32 ptrsem;                  /* Receiver semaphore          */
    uint16 ptstate;                /* Port state (FREE/LIMBO/ALLOC) */
    uint16 ptmaxcnt;               /* Max messages to be queued   */
    int32 ptseq;                   /* Sequence changed at creation */
    struct ptnode *pthead;         /* List of message pointers    */
    struct ptnode *pttail;        /* Tail of message list        */
};

extern struct ptnode *ptfree;      /* List of free nodes          */
extern struct ptentry porttab[];   /* Port table                   */
extern int32 ptnextid;            /* Next port ID to try when    */
                                  /* looking for a free slot     */

#define isbadport(portid) ( (portid)<0 || (portid)>=NPORTS )
```

Xinu Port Functions

- *Ptinit*
 - Called once at startup
 - Initializes port system
- *Ptcreate*
 - Creates a new port
 - Argument specifies number of messages
- *Ptsend*
 - Sends a message to a port
- *Ptrecv*
 - Retrieves a message from a port

Xinu Port Functions

(continued)

- *Ptreset*
 - Resets existing port
 - Disposes of existing messages
 - Allows waiting processes to continue
- *Ptdelete*
 - Deletes existing port
 - Disposes of existing messages
 - Allows blocked processes to continue

Xinu Ptinit (part 1)

```
/* ptinit.c - ptinit */
```

```
#include <xinu.h>
```

```
struct ptnode *ptfree;  
struct ptnode porttab[NPORTS];  
int32 ptnextid;
```

1. Make each port free
2. Form the linked list of free nodes
3. Initialize *ptnextid*

```
/* List of free message nodes */  
/* Port table */  
/* Next table entry to try */
```

```
/*-----  
 * ptinit - Initialize all ports  
 *-----  
 */  
syscall ptinit(  
    int32 maxmsgs /* Total messages in all ports */  
)  
{  
    int32 i; /* Runs through the port table */  
    struct ptnode *next, *curr; /* Used to build a free list */  
  
    /* Allocate memory for all messages on all ports */  
    ptfree = (struct ptnode *)getmem(maxmsgs*sizeof(struct ptnode));  
    if (ptfree == (struct ptnode *)SYSERR) {  
        panic("pinit - insufficient memory");  
    }  
}
```

allocate a block
of memory

Xinu Ptinit (part 2)

```
/* Initialize all port table entries to free */
```

```
for (i=0 ; i<NPORTS ; i++) {  
    porttab[i].ptstate = PT_FREE;  
    porttab[i].ptseq = 0;
```

The search will start when a new port is needed.

```
}  
ptnextid = 0;
```

Give the index in array

```
/* Create a free list of message nodes linked together */
```

```
for ( curr=next=ptfree ; --maxmsgs > 0 ; curr=next ) {  
    curr->ptnext = ++next;  
}
```

```
/* Set the pointer in the final node to NULL */
```

```
curr->ptnext = NULL;
```

```
return OK;
```

```
}
```

Xinu Ptcreate (part 1)

```
/* ptcreate.c - ptcreate */
```

```
#include <xinu.h>
```

Specify the maximum count of outstanding messages that the port will allow

```
/*-----  
 * ptcreate - Create a port that allows "count" outstanding messages  
 *-----  
 */  
syscall ptcreate(  
    int32 count                /* Size of port */  
)  
{  
    intmask mask;              /* Saved interrupt mask */  
    int32 i;                   /* Counts all possible ports */  
    int32 ptnum;               /* Candidate port number to try */  
    struct ptentry *ptptr;     /* Pointer to port table entry */  
  
    mask = disable();  
    if (count < 0) {  
        restore(mask);  
        return SYSERR;  
    }  
}
```

Xinu Ptcreate (part 2)

```
for (i=0 ; i<NPORTS ; i++) {    /* Count all table entries    */
    ptnum = ptnextid;           /* Get an entry to check    */
    if (++ptnextid >= NPORTS) {
        ptnextid = 0;          /* Reset for next iteration */
    }

    /* Check table entry that corresponds to ID ptnum */

    ptptr= &porttab[ptnum];
    if (ptptr->ptstate == PT_FREE) {
        ptptr->ptstate = PT_ALLOC;
        ptptr->ptssem = semcreate(count);
        ptptr->ptrsem = semcreate(0);
        ptptr->pthead = ptptr->pttail = NULL;
        ptptr->ptseq++;
        ptptr->ptmaxcnt = count;
        restore(mask);
        return ptnum;
    }
}
restore(mask);
return SYSERR;
}
```

Allocate an entry in the port table from

Return a port identifier (port ID) to its caller

Xinu Ptsend (part 1)

```
/* ptsend.c - ptsend */
```

```
#include <xinu.h>
```

1. Enqueue the message
2. Signal the receiver semaphore
3. Return

```
/*-----  
 * ptsend - Send a message to a port by adding it to the waiting queue  
 *-----
```

```
 */  
syscall ptsend(  
    int32 portid,          /* ID of port to use          */  
    umsg32 msg             /* Message to send          */  
)
```

Passing a port ID as an argument

```
{  
    intmask mask;          /* Saved interrupt mask      */  
    struct ptentry *ptptr; /* Pointer to table entry    */  
    int32 seq;             /* Local copy of sequence num.*/  
    struct ptnode *msgnode; /* Allocated message node    */  
    struct ptnode *tailnode; /* Last node in port or NULL */  
}
```

```
mask = disable();  
if (isbadport(portid) ||  
    (ptptr = &porttab[portid])->ptstate != PT_ALLOC ) {  
    restore(mask);  
    return SYSERR;  
}
```

Specify a valid port ID

Xinu Ptsend (part 2)

```
/* wait for space and verify port has not been reset */
```

A local copy of the sequence number

```
seq = ptptr->ptseq; /* Record original sequence */
```

```
if (wait(ptptr->ptssem) == SYSERR  
    || ptptr->ptstate != PT_ALLOC  
    || ptptr->ptseq != seq) {  
    restore(mask);  
    return SYSERR;  
}
```

1. Wait sender semaphore
2. Verify that the port is still allocated
3. The sequence number agrees.

```
if (ptfree == NULL) {  
    panic("Port system ran out of message nodes");  
}
```

```
/* Obtain node from free list by unlinking */
```

```
msgnode = ptfree; /* Point to first free node */  
ptfree = msgnode->ptnext; /* Unlink from the free list */  
msgnode->ptnext = NULL; /* Set fields in the node */  
msgnode->ptmsg = msg;
```

Xinu Ptsend (part 3)

```
/* Link into queue for the specified port */
```

```
tailnode = ptptr->pttail;
```

Enqueues messages in FIFO order

```
if (tailnode == NULL) {
```

```
/* Queue for port was empty
```

```
*/
```

```
    ptptr->pttail = ptptr->pthead = msgnode;
```

```
} else {
```

```
/* Insert new node at tail
```

```
*/
```

```
    tailnode->ptnext = msgnode;
```

```
    ptptr->pttail = msgnode;
```

```
}
```

```
signal(ptptr->ptrsem);
```

```
restore(mask);
```

```
return OK;
```

```
}
```

Point to a new node after the node has been added to the list

Signal the receiver semaphore after added a new message

Xinu Ptrecev (part 1)

```
/* ptrecev.c - ptrecev */
```

```
#include <xinu.h>
```

1. Remove a message from a specified port
2. Return the message to its caller

```
/*-----  
 * ptrecev - Receive a message from a port, blocking if port empty  
 *-----  
 */  
uint32 ptrecev(  
    int32 portid                /* ID of port to use          */  
)  
{  
    intmask mask;               /* Saved interrupt mask      */  
    struct ptenry *ptptr;       /* Pointer to table entry    */  
    int32 seq;                  /* Local copy of sequence num. */  
    umsg32 msg;                 /* Message to return         */  
    struct ptnode *msgnode;     /* First node on message list */  
  
    mask = disable();  
    if ( isbadport(portid) ||  
        (ptptr= &porttab[portid])->ptstate != PT_ALLOC ) {  
        restore(mask);  
        return (uint32)SYSERR;  
    }  
}
```

Xinu Ptrecev (part 2)

```
/* wait for message and verify that the port is still allocated */
```

```
seq = ptptr->ptseq; /* Record original sequence */
if (wait(ptptr->ptrsem) == SYSERR
    || ptptr->ptstate != PT_ALLOC
    || ptptr->ptseq != seq) {
    restore(mask);
    return (uint32)SYSERR;
}
```

1. Wait until a msg is available
2. Verify that the port has not been deleted or reused

```
/* Dequeue first message that is waiting in the port */
```

Record the value of the message

```
msgnode = ptptr->pthead;
msg = msgnode->ptmsg;
if (ptptr->pthead == ptptr->pttail) /* Delete last item */
    ptptr->pthead = ptptr->pttail = NULL;
else
    ptptr->pthead = msgnode->ptnext;
msgnode->ptnext = ptfree; /* Return to free list */
ptfree = msgnode;
signal(ptptr->ptssem);
restore(mask);
return msg;
```

```
}
```

Port Reset and Deletion

- *Ptreset* or *ptdelete*
 - Disposes of existing messages, if the port contains any
 - Unblocks processes that are waiting
 - To send
 - To receive
- Semaphores are either reset or deleted
- Processes are informed that an abnormal termination occurred

How should the port system dispose of waiting messages?

Disposing of Messages

- Needed during reset and deletion
- Alternatives
 - Fixed set of choices
 - Allow arbitrary processing
- Arbitrary processing: Permit a caller to specify how to dispose of messages
 - More general
 - Must allow caller to specify disposition function as argument to *ptreset* or *ptdelete*
 - Disposition function is called for each existing message

Xinu Ptdelete

```
/* ptdelete.c - ptdelete */
```

```
#include <xinu.h>
```

```
/*-----  
 * ptdelete - Delete a port, freeing waiting processes and messages  
 *-----  
 */  
syscall ptdelete(  
    int32 portid,                /* ID of port to delete */  
    int32 (*disp)(int32)        /* Function to call to dispose  
                                /* of waiting messages */  
)  
{  
    intmask mask;                /* Saved interrupt mask */  
    struct ptenry *ptptr;        /* Pointer to port table entry */  
  
    mask = disable();  
    if ( isbadport(portid) ||  
        (ptptr= &porttab[portid])->ptstate != PT_ALLOC ) {  
        restore(mask);  
        return SYSERR;  
    }  
    _ptclear(ptptr, PT_FREE, disp);  
    ptnextid = portid;  
    restore(mask);  
    return OK;  
}
```

- Perform the work of clearing messages and waiting processes
- Place the port in a “limbo” state (PT_LIMBO).

Xinu Ptrreset

```
/* ptrreset.c - ptrreset */
```

```
#include <xinu.h>
```

```
/*-----  
 * ptrreset - Reset a port, freeing waiting processes and messages and  
 * leaving the port ready for further use  
 *-----  
 */  
syscall ptrreset(  
    int32 portid,          /* ID of port to reset          */  
    int32 (*disp)(int32)  /* Function to call to dispose */  
    )                    /* of waiting messages        */  
{  
    intmask mask;          /* saved interrupt mask        */  
    struct ptenry *ptptr;  /* Pointer to port table entry */  
  
    mask = disable();  
    if ( isbadport(portid) ||  
        (ptptr= &porttab[portid])->ptstate != PT_ALLOC ) {  
        restore(mask);  
        return SYSERR;  
    }  
    _ptclear(ptptr, PT_ALLOC, disp);  
    restore(mask);  
    return OK;  
}
```


Xinu _ptclear (part 1)

```
/* ptclear.c - _ptclear */
```

```
#include <xinu.h>
```

```
/*-----  
 * _ptclear - Used by ptdelete and ptreset to clear or reset a port  
 * (internal function assumes interrupts disabled and  
 * arguments have been checked for validity)  
 *-----  
 */  
void _ptclear(  
    struct ptentry *ptptr,          /* Table entry to clear          */  
    uint16 newstate,               /* New state for port            */  
    int32 (*dispose)(int32)        /* Disposal function to call     */  
)  
{  
    struct ptnode *walk; /* Pointer to walk message list */  
  
    /* Place port in limbo state while waiting processes are freed */  
  
    ptptr->ptstate = PT_LIMBO;  
  
    ptptr->ptseq++; /* Reset accession number */  
    walk = ptptr->pthead; /* First item on msg list */  
}
```

Waiting processes can tell that the
port has changed when they awaken

Xinu _ptclear (part 2)

```
if ( walk != NULL ) {                                /* If message list nonempty */

    /* walk message list and dispose of each message */
    for( ; walk!=NULL ; walk=walk->ptnext) {
        (*dispose)( walk->ptmsg );
    }

    /* Link entire message list into the free list */

    (ptptr->pttail)->ptnext = ptfree;
    ptfree = ptptr->pthead;
}

if (newstate == PT_ALLOC) {
    ptptr->pttail = ptptr->pthead = NULL;
    semreset(ptptr->ptssem, ptptr->ptmaxcnt);
    semreset(ptptr->ptrsem, 0);
} else {
    semdelete(ptptr->ptssem);
    semdelete(ptptr->ptrsem);
}
ptptr->ptstate = newstate;
return;
}
```

Delete or reset the semaphores given by its second argument

Concurrency and Message Disposition

- Disposition routine
 - Is specified by user
 - May reschedule
- Example
 - Message to be deleted contains a pointer to a buffer
 - Disposition routine calls *freebuf* to release a buffer back to the pool.

Consequence: concurrency problems may arise

Semaphore Reset and Deletion

- Resetting or deleting a port will reset or delete the semaphores
- If processes are blocked on the semaphore, they will become ready
- The rescheduling invariant means a higher priority process may execute
- Additional processes may attempt to use the port

Consequence: we need to handle attempts to use the port concurrently during reset or deletion

Three Mechanisms for Handling Reset

(1/3)

- Accession numbers: assign a sequence number to a port
 - Increment sequence
 - When port created and
 - When port deleted or reset
 - Have *ptsend* and *ptrecv* record sequence number when operation begins and check sequence number after *wait* returns
 - If sequence number changed, port was reset, so operation should abort

Three Mechanisms for Handling Reset

(2/3)

- New state for the port: assign each port a *state* variable
 - Values of the state variable
 - *PTFREE* if not in use
 - *PTALLOC* if in use
 - *PTLIMBO* if in transition
 - Have *ptsend* and *ptrecv* examine state variable
 - If state is *PTLIMBO* port is being reset or deleted and cannot be used

Three Mechanisms for Handling Reset

(3/3)

- Deferred rescheduling: temporarily postpone scheduling decisions
 - Call *resched_cntl(DEFER_START)* at start of reset or delete
 - Call *resched_cntl(DEFER_START)* after all operations are performed

Summary (1/2)

- Xinu offers a low-level message passing mechanism
 - Process-to-process
 - Only one message outstanding
- Xinu offers a high-level message passing mechanism
 - Dynamically created ports
 - Number of messages and message size fixed when port created
 - Arbitrary senders and receivers
 - Synchronous interface

Summary (2/2)

- Port reset /deletion is tricky because
 - Unblocked processes may execute
 - New processes may attempt to use the port
 - Three techniques can handle transition
 - Sequence number informs waiting processes that port is being reset or deleted
 - Limbo state prevents new processes from using the port while it is being reset or deleted
 - Deferred rescheduling